

ACT1_Regresion

September 10, 2024

1 Actividad 1: Problemas de Regresión

A01285158 | Grace Aviance Silva Aróstegui

2 EJERCICIO 1

El conjunto de datos de esperanzas de vida (Life Expectancy (WHO) | Kaggle) tiene el registro de la esperanza de vida de 193 países medida en diferentes años, junto con otras variables que se pueden relacionar con riesgos a la salud y la mortalidad.

Para este ejercicio, sólo se considerará como variable dependiente la cuarta columna (“Life expectancy”). A su vez, las variables independientes de interés son:

- X1 - Adult mortality
- X2 - Infant deaths
- X3 - Alcohol
- X4 - Percentage expenditure
- X5 - Hepatitis B
- X6 - Measles
- X7 - BMI
- X8 - Under-five deaths
- X9 - Polio
- X10 - Total expenditure
- X11 - Diphtheria
- X12 - HIV/AIDS
- X13 - GDP
- X14 - Population
- X15 - Thinness 1-19 years
- X16 - Thinness 5-9 years
- X17 - Income composition of resources
- X18 - Schooling

Nota 1: Las variables con las que vas a trabajar depende del **penúltimo número de tu matrícula** de acuerdo a la siguiente lista:

- 0, 1 - Todas las variables, menos X1, X5, X9, X13, X17
- 2, 3 - Todas las variables, menos X2, X6, X10, X14, X18
- **4, 5 - Todas las variables, menos X3, X7, X11, X15**
- 6, 7 - Todas las variables, menos X4, X8, X12, X16

- 8, 9 - Todas las variables, menos X5, X9, X13, X17

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.preprocessing import StandardScaler
import warnings
warnings.filterwarnings("ignore")
```

```
[2]: # Primeramente reduciremos el dataset manteniendo las que sean de interes
df = pd.read_csv("life_expectancy_data.csv")
df.columns
```

```
[2]: Index(['Country', 'Year', 'Status', 'Life expectancy ', 'Adult Mortality',
'infant deaths', 'Alcohol', 'percentage expenditure', 'Hepatitis B',
'Measles ', ' BMI ', 'under-five deaths ', 'Polio', 'Total expenditure',
'Diphtheria ', ' HIV/AIDS', 'GDP', 'Population',
' thinness 1-19 years', ' thinness 5-9 years',
'Income composition of resources', 'Schooling'],
dtype='object')
```

```
[3]: extra_columns = ['Country', 'Year', 'Status']
df = df.drop(extra_columns, axis=1)
df.columns
```

```
[3]: Index(['Life expectancy ', 'Adult Mortality', 'infant deaths', 'Alcohol',
'percentage expenditure', 'Hepatitis B', 'Measles ', ' BMI ',
'under-five deaths ', 'Polio', 'Total expenditure', 'Diphtheria ',
' HIV/AIDS', 'GDP', 'Population', ' thinness 1-19 years',
' thinness 5-9 years', 'Income composition of resources', 'Schooling'],
dtype='object')
```

```
[4]: # Excluyendo aquellas de acuerdo a la "Nota 1", las cuales son todas menos X3_
↳ (Alcohol), X7 (BMI), X11 (Diphtheria), X15 (thinness 5-9 years)
columns_note1 = ['Alcohol', ' BMI ', 'Diphtheria ', ' thinness 1-19 years']
df = df.drop(columns_note1, axis=1)
df.columns
```

```
[4]: Index(['Life expectancy ', 'Adult Mortality', 'infant deaths',
'percentage expenditure', 'Hepatitis B', 'Measles ',
'under-five deaths ', 'Polio', 'Total expenditure', ' HIV/AIDS', 'GDP',
'Population', ' thinness 5-9 years', 'Income composition of resources',
'Schooling'],
dtype='object')
```

```
[5]: # Ya tenemos nuestro Dataset. Revisaremos que no hayan datos Nulos
df.isnull().sum()
```

```
[5]: Life expectancy          10
Adult Mortality              10
infant deaths                0
percentage expenditure       0
Hepatitis B                  553
Measles                      0
under-five deaths            0
Polio                        19
Total expenditure            226
HIV/AIDS                     0
GDP                          448
Population                   652
  thinness 5-9 years          34
Income composition of resources 167
Schooling                    163
dtype: int64
```

```
[6]: df.dropna(subset=['Life expectancy ', 'Adult Mortality', 'Hepatitis B',
↳ 'Polio', 'Total expenditure',
      'GDP', 'Population', ' thinness 5-9 years', 'Income
↳ composition of resources',
      'Schooling'], inplace=True)
df.isnull().sum()
```

```
[6]: Life expectancy          0
Adult Mortality              0
infant deaths                0
percentage expenditure       0
Hepatitis B                  0
Measles                      0
under-five deaths            0
Polio                        0
Total expenditure            0
HIV/AIDS                     0
GDP                          0
Population                   0
  thinness 5-9 years          0
Income composition of resources 0
Schooling                    0
dtype: int64
```

1. Grafica cada variable predictora vs la variable de respuesta asignadas a tu número de matrícula.

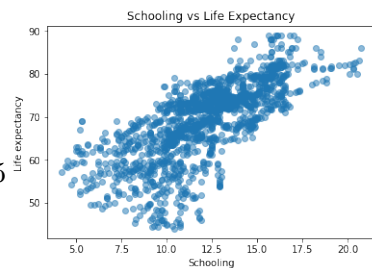
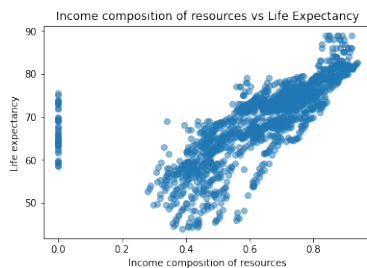
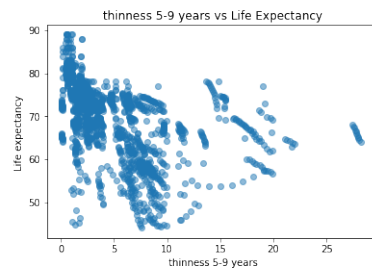
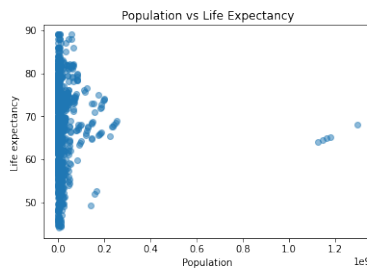
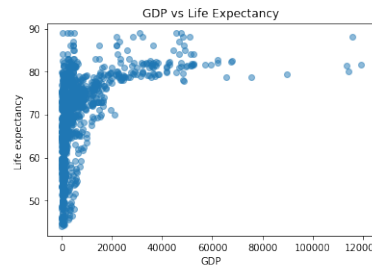
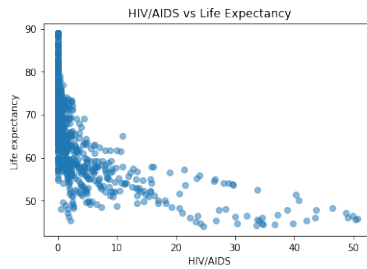
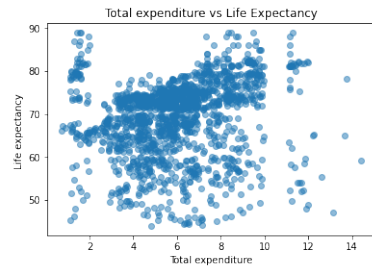
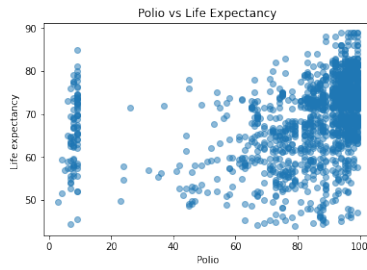
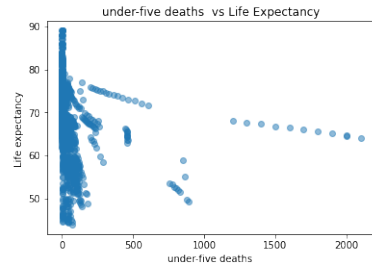
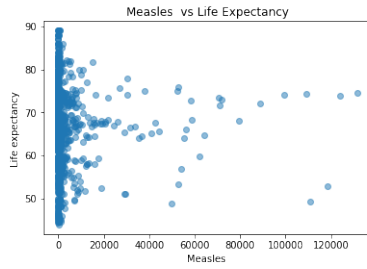
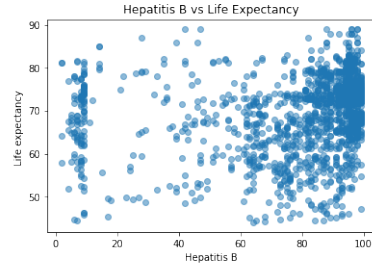
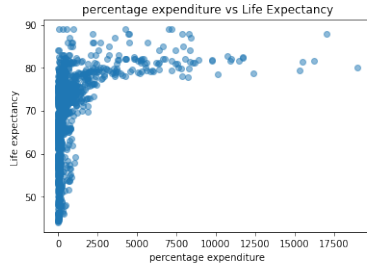
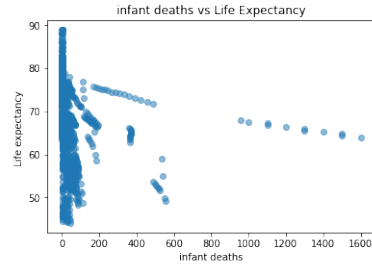
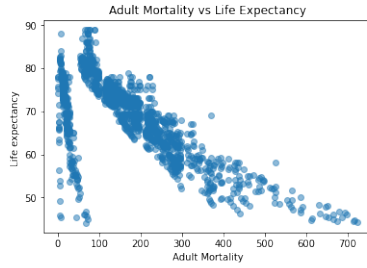
```
[7]: # Variables predictorias
vars_independientes = ['Adult Mortality', 'infant deaths', 'percentage_
↳expenditure', 'Hepatitis B', 'Measles ',
                        'under-five deaths ', 'Polio', 'Total expenditure', ' HIV/
↳AIDS', 'GDP',
                        'Population', ' thinness 5-9 years', 'Income composition of_
↳resources', 'Schooling']

fig, axs = plt.subplots(7, 2, figsize=(12, 28)) # 7 renglones, 2 columnas,
↳tamaño de la figura
fig.tight_layout(pad=4.0) # Espacio entre gráficos

for i, var in enumerate(vars_independientes):
    row = i // 2 # Determina el índice de la fila
    col = i % 2  # Determina el índice de la columna

    axs[row, col].scatter(df[var], df['Life expectancy '], alpha=0.5)
    axs[row, col].set_xlabel(var)
    axs[row, col].set_ylabel('Life expectancy ')
    axs[row, col].set_title(f'{var} vs Life Expectancy')

plt.show()
```



2. Implementa la fórmula directa para calcular los coeficientes de un modelo de regresión lineal, y obtenga con ella el modelo que corresponde a la variable de respuesta y las variables predictoras asignadas a tu número de matrícula. Dado que tenemos varias variables predictoras (independientes), nuestro modelo de Regresión Lineal sería Múltiple. Existe la Regresión Lineal Múltiple con optimización de Descenso Gradiente. Elegí sin descenso gradiente ya que: - El tamaño del dataset es pequeño, tiene menos de 10,000 muestras y menos de 100 características - No hay problemas de multicolinealidad

```
[8]: # Antes de hacer el modelo, escalamos los datos

x = df[['Adult Mortality', 'infant deaths', 'percentage expenditure',
        ↪ 'Hepatitis B', 'Measles ',
        ↪ 'under-five deaths ', 'Polio', 'Total expenditure', ' HIV/
        ↪ AIDS', 'GDP',
        ↪ 'Population', ' thinness 5-9 years', 'Income composition of
        ↪ resources', 'Schooling']]
y = df['Life expectancy '] # Variable de respuesta

x_scaled = StandardScaler().fit_transform(x)

# Matriz con valores de las variables independientes escaladas
X = np.column_stack((np.ones(x_scaled.shape[0]), x_scaled))

# Fit model function (it fits a linear model using the specified data set).
def fit_model(X, y):
    return np.linalg.inv(X.transpose() @ X) @ X.transpose() @ y

# Predict function (it evaluates an array of observations using the specified
    ↪ linear model).
def predict(X, beta):
    return X @ beta

# Build linear model
beta = fit_model(X, y)
#np.set_printoptions(suppress=True)
print ("Model coefficients: ", beta)
```

```
Model coefficients: [ 6.93023044e+01 -2.18560185e+00  1.22898976e+01
 6.77917958e-01
 -9.97287431e-03 -1.47856294e-01 -1.24489688e+01  2.53510851e-01
 1.95191864e-01 -2.67056877e+00  8.99806028e-02 -1.87859898e-02
 -4.87343711e-01  1.86699223e+00  2.55796319e+00]
```

3. Evalúa con validación cruzada de k-pliegues tu modelo, calculando los valores de R^2 , MSE y MAE

```
[9]: # Residuals
y_pred = predict(X, beta)
r = y - y_pred

# Calculate MSE, MAE and R^2 with the training set
print('MSE: ', mean_squared_error(y, y_pred))
print("MAE: ", mean_absolute_error(y, y_pred))
print("R^2: ", r2_score(y, y_pred))
```

```
MSE: 13.08532811014522
MAE: 2.773793947439496
R^2: 0.8308019870625327
```

```
[10]: X = np.column_stack((np.ones(x_scaled.shape[0]), x_scaled))
y = y.values
# Evaluate model with cross validation
n_folds = 5
kf = KFold(n_splits=n_folds, shuffle = True)

mse_cv = []
mae_cv = []
r2_cv = []
for train_index, test_index in kf.split(X):

    # Training phase
    x_train = X[train_index, :]
    y_train = y[train_index]

    beta_cv = fit_model(x_train, y_train)

    # Test phase
    x_test = X[test_index, :]
    y_test = y[test_index]

    y_pred = predict(x_test, beta_cv)

    # Calculate MSE, MAE and R^2
    mse_i = mean_squared_error(y_test, y_pred)
    #print('mse = ', mse_i)
    mse_cv.append(mse_i)

    mae_i = mean_absolute_error(y_test, y_pred)
    #print('mae = ', mae_i)
    mae_cv.append(mae_i)

    r2_i = r2_score(y_test, y_pred)
    #print('r^2= ', r2_i)
    r2_cv.append(r2_i)
```

```
print('MSE:', np.average(mse_cv), ' MAE:', np.average(mae_cv), ' R^2:', np.
      ↪average(r2_cv))
```

MSE: 13.391850356756162 MAE: 2.8048269248638156 R^2: 0.8261617260877581

4. Utiliza validación cruzada de Monte Carlo con 1000 iteraciones para encontrar histogramas de R^2 , MSE y MAE.

```
[13]: from sklearn.model_selection import train_test_split

# Listas para almacenar los resultados de cada iteración
mse_list = []
mae_list = []
r2_list = []

# Monte Carlo, 1000 iteraciones
for _ in range(1000):
    # Dividir los datos. 80% Entrenamiento, 20% Prueba
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

    # Ajustar el modelo
    beta = fit_model(X_train, y_train)

    # Hacer predicciones
    y_pred = predict(X_test, beta)

    # Calcular MSE, MAE y R^2
    mse_list.append(mean_squared_error(y_test, y_pred))
    mae_list.append(mean_absolute_error(y_test, y_pred))
    r2_list.append(r2_score(y_test, y_pred))

print('MSE:', np.average(mse_list), ' MAE:', np.average(mae_list), ' R^2:', np.
      ↪average(r2_list))

# Graficar los histogramas
plt.figure(figsize=(15, 5))

plt.subplot(1, 3, 1)
plt.hist(mse_list, bins=30, color='skyblue', edgecolor='black')
plt.title('Histogram of MSE')
plt.xlabel('MSE')
plt.ylabel('Frequency')

plt.subplot(1, 3, 2)
plt.hist(mae_list, bins=30, color='lightgreen', edgecolor='black')
plt.title('Histogram of MAE')
plt.xlabel('MAE')
```

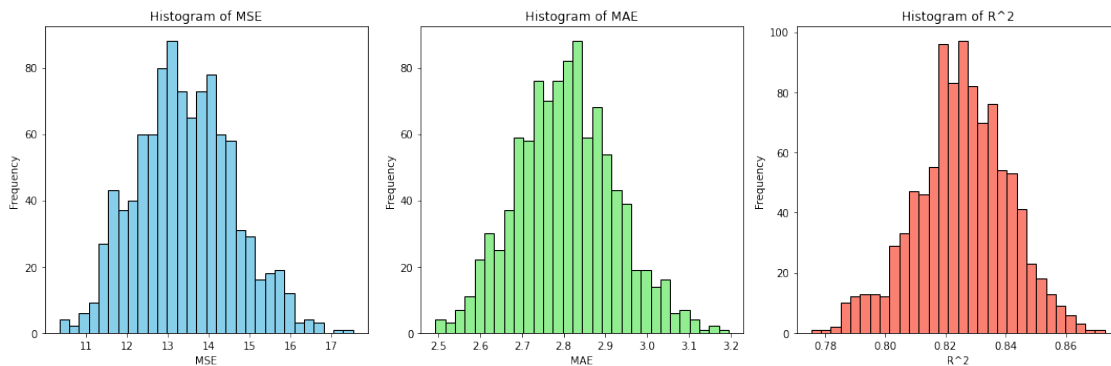


```
plt.ylabel('Frequency')

plt.subplot(1, 3, 3)
plt.hist(r2_list, bins=30, color='salmon', edgecolor='black')
plt.title('Histogram of R^2')
plt.xlabel('R^2')
plt.ylabel('Frequency')

plt.tight_layout()
plt.show()
```

MSE: 13.422741567779054 MAE: 2.8047238405132267 R^2: 0.825507725651008



5. Utiliza el método de validación cruzada asignado a tu matrícula para mostrar los histogramas de MSE y MAE. ¿Los histogramas son distintos a los obtenidos con el método de Monte Carlo?

```
[14]: # Evaluate model with LOOCV
n_samples = len(X)
kf = KFold(n_splits=n_samples, shuffle=True)

mse_cv = []
mae_cv = []
for train_index, test_index in kf.split(X):

    # Training phase
    x_train = X[train_index, :]
    y_train = y[train_index]
    beta_cv = fit_model(x_train, y_train)

    # Test phase
    x_test = X[test_index, :]
    y_test = y[test_index]
    y_pred = predict(x_test, beta_cv)
```

```

# Calculate MSE, MAE, and R^2
mse_i = mean_squared_error(y_test, y_pred)
mse_cv.append(mse_i)

mae_i = mean_absolute_error(y_test, y_pred)
mae_cv.append(mae_i)

# Summary statistics
print('MSE:', np.average(mse_cv), ' MAE:', np.average(mae_cv))

plt.figure(figsize=(18, 6))
plt.subplot(1, 3, 1)

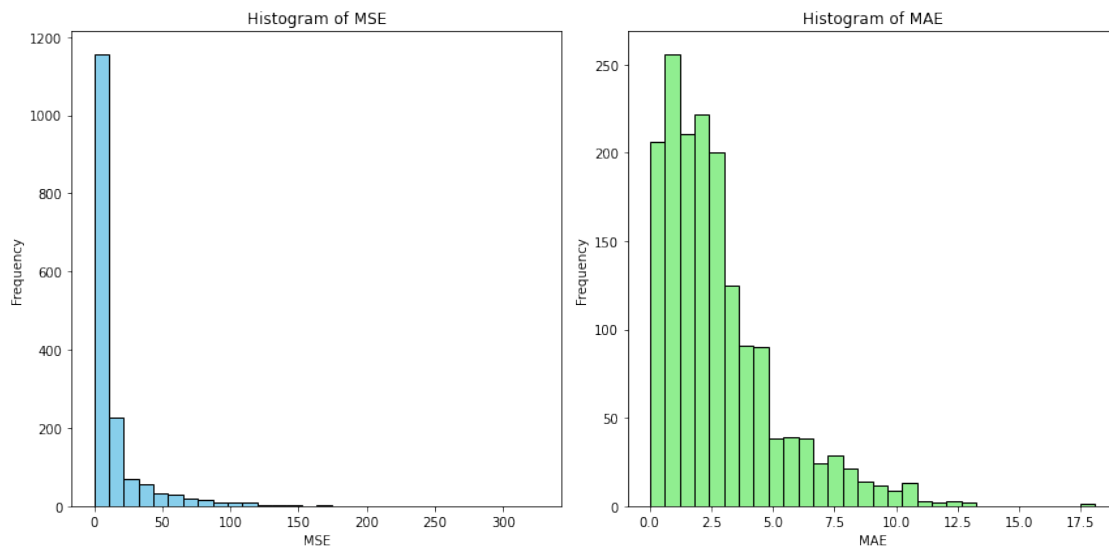
plt.subplot(1, 3, 1)
plt.hist(mse_cv, bins=30, color='skyblue', edgecolor='black')
plt.title('Histogram of MSE')
plt.xlabel('MSE')
plt.ylabel('Frequency')

plt.subplot(1, 3, 2)
plt.hist(mae_cv, bins=30, color='lightgreen', edgecolor='black')
plt.title('Histogram of MAE')
plt.xlabel('MAE')
plt.ylabel('Frequency')

plt.tight_layout()
plt.show()

```

MSE: 13.374834581890257 MAE: 2.8018956192333357



La diferencia entre los histogramas de validación cruzada y Monte Carlo son totalmente diferentes. Los histogramas de Monte Carlo tiene mayor variabilidad (lo cual tiene sentido por el hecho de ser un método mucho más robusto) y tiende a una distribución similar a la gaussiana. Mientras que los histogramas de validación cruzada están muy sesgadas a la derecha, con la mayoría de los datos tendiendo a agruparse a las primeras clases

6. Agrega al conjunto de datos columnas que representen los cuadrados de las variables predictoras (por ejemplo, X_{11}^2 , X_{13}^2), así como los productos entre pares de variables (por ejemplo, $X_1 \times X_2$, $X_3 \times X_4$). Repita los pasos 1, 2 y 3 pero con este nuevo conjunto de datos.

```
[20]: df.columns
```

```
[20]: Index(['Life expectancy ', 'Adult Mortality', 'infant deaths',
          'percentage expenditure', 'Hepatitis B', 'Measles ',
          'under-five deaths ', 'Polio', 'Total expenditure', ' HIV/AIDS', 'GDP',
          'Population', ' thinness 5-9 years', 'Income composition of resources',
          'Schooling'],
          dtype='object')
```

```
[21]: df['Polio_x_HepatitisB'] = df['Polio'] * df['Hepatitis B']
df['GDP_x_Population'] = df['GDP'] * df['Population']
df['Schooling_x_Measles'] = df['Schooling'] * df['Measles ']

df['Polio_^2'] = df['Polio'] ** 2
df['Infant_Deaths_^2'] = df['infant deaths'] ** 2
df['Total_expenditure_^2'] = df['Total expenditure'] ** 2

df6 = df
df6.head(4)
```

```
[21]:
```

	Life expectancy	Adult Mortality	infant deaths	percentage expenditure \
0	65.0	263.0	62	71.279624
1	59.9	271.0	64	73.523582
2	59.9	268.0	66	73.219243
3	59.5	272.0	69	78.184215

	Hepatitis B	Measles	under-five deaths	Polio	Total expenditure \
0	65.0	1154	83	6.0	8.16
1	62.0	492	86	58.0	8.18
2	64.0	430	89	62.0	8.13
3	67.0	2787	93	67.0	8.52

	HIV/AIDS	...	Population	thinness 5-9 years \
0	0.1	...	33736494.0	17.3
1	0.1	...	327582.0	17.5
2	0.1	...	31731688.0	17.7
3	0.1	...	3696958.0	18.0

	Income composition of resources	Schooling	Polio_x_HepatitisB	\
0	0.479	10.1	390.0	
1	0.476	10.0	3596.0	
2	0.470	9.9	3968.0	
3	0.463	9.8	4489.0	

	GDP_x_Population	Schooling_x_Measles	Polio_^2	Infant_Deaths_^2	\
0	1.971086e+10	11655.4	36.0	3844	
1	2.007083e+08	4920.0	3364.0	4096	
2	2.004633e+10	4257.0	3844.0	4356	
3	2.476810e+09	27312.6	4489.0	4761	

	Total_expenditure_^2
0	66.5856
1	66.9124
2	66.0969
3	72.5904

[4 rows x 21 columns]

```
[44]: df6.columns
```

```
[44]: Index(['Life expectancy ', 'Adult Mortality', 'infant deaths',
        'percentage expenditure', 'Hepatitis B', 'Measles ',
        'under-five deaths ', 'Polio', 'Total expenditure', ' HIV/AIDS', 'GDP',
        'Population', ' thinness 5-9 years', 'Income composition of resources',
        'Schooling', 'Polio_x_HepatitisB', 'GDP_x_Population',
        'Schooling_x_Measles', 'Polio_^2', 'Infant_Deaths_^2',
        'Total_expenditure_^2'],
        dtype='object')
```

```
[45]: ### PUNTO 1
### 1. Grafica cada variable predictora vs
### la variable de respuesta asignadas a tu número de matrícula.

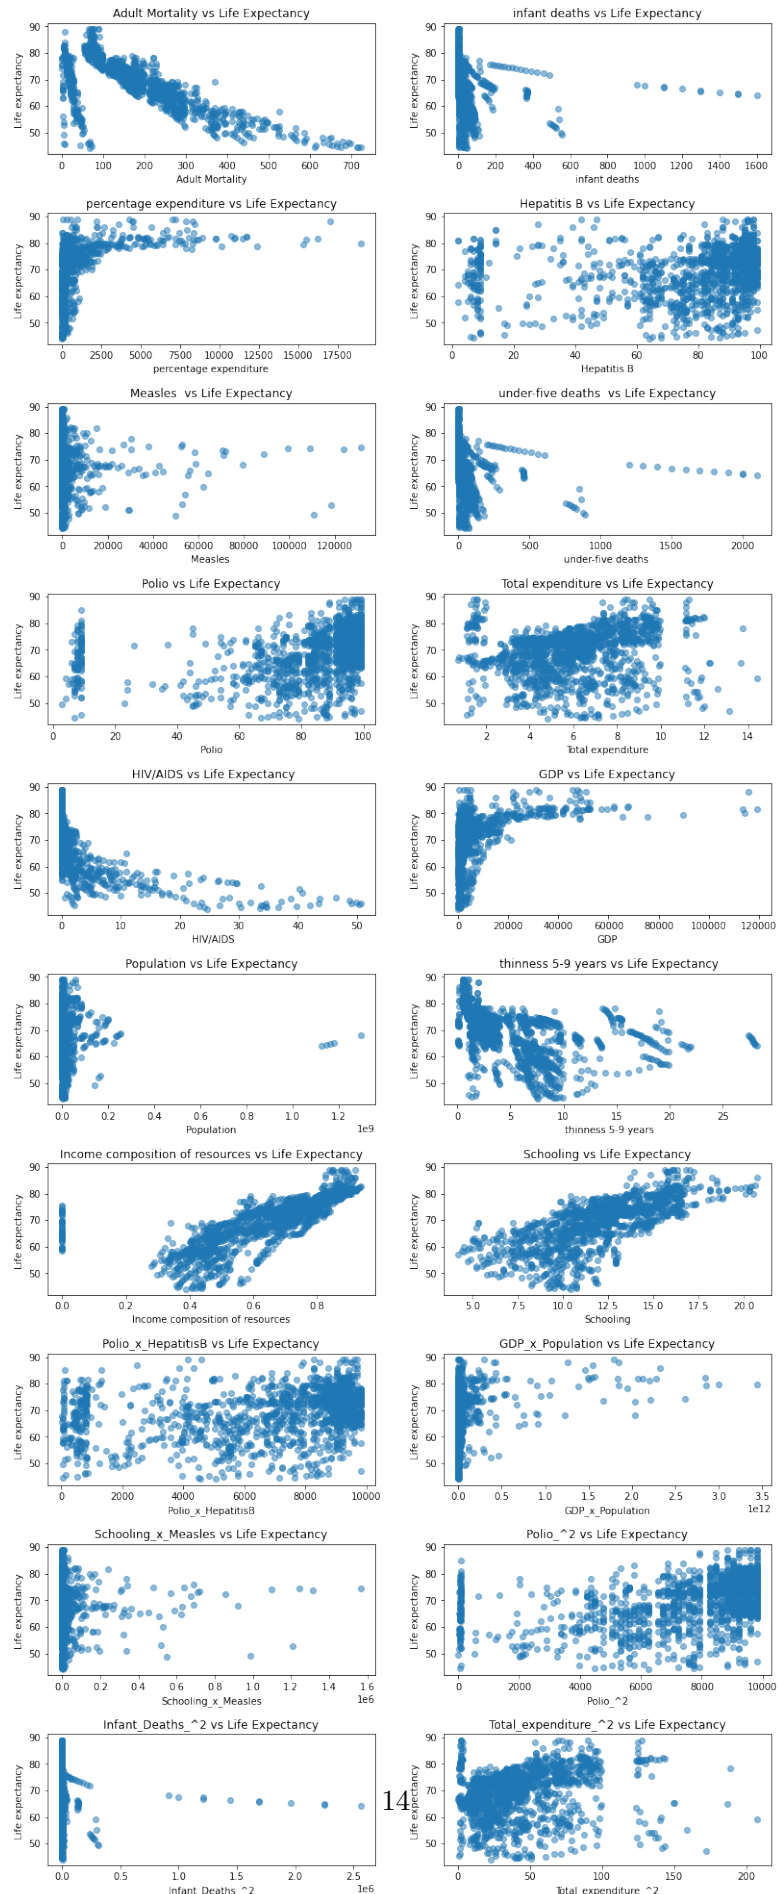
# Variables predictoras
vars_independientes_6 = ['Adult Mortality', 'infant deaths',
        'percentage expenditure', 'Hepatitis B', 'Measles ',
        'under-five deaths ', 'Polio', 'Total expenditure', ' HIV/AIDS', 'GDP',
        'Population', ' thinness 5-9 years', 'Income composition of resources',
        'Schooling', 'Polio_x_HepatitisB', 'GDP_x_Population',
        'Schooling_x_Measles', 'Polio_^2', 'Infant_Deaths_^2',
        'Total_expenditure_^2']

fig, axs = plt.subplots(10, 2, figsize=(12, 28)) # 10 renglones, 2 columnas
fig.tight_layout(pad=4.0) # Espacio entre gráficos
```

```
for i, var in enumerate(vars_independientes_6):
    row = i // 2 # Determina el índice de la fila
    col = i % 2  # Determina el índice de la columna

    axs[row, col].scatter(df6[var], df6['Life expectancy '], alpha=0.5)
    axs[row, col].set_xlabel(var)
    axs[row, col].set_ylabel('Life expectancy ')
    axs[row, col].set_title(f'{var} vs Life Expectancy')

plt.show()
```



```
[48]: df6['Adult Mortality'].dtype
```

```
[48]: dtype('float64')
```

```
[52]: ### PUNTO 2
      ### Obtener los coeficientes del modelo de regresion

      import numpy as np
      import matplotlib.pyplot as plt

      from sklearn.model_selection import KFold
      from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

      # Fit model function (it fits a linear model using the specified data set).
      def fit_model(X_6, y_6):
          return np.linalg.inv(X_6.transpose() @ X_6) @ X_6.transpose() @ y_6

      # Predict function (it evaluates an array of observations using the specified
      ↪ linear model).
      def predict(X_6, beta_6):
          return X_6 @ beta_6

      # Generate data
      x_6 = df6.values
      X_6 = np.column_stack((np.ones(x_6.shape[0]), x_6))

      # Variable de respuesta
      y_6 = df6['Life expectancy '].values

      # Build linear model
      beta_6 = fit_model(X_6, y_6)
      print ("Model coefficients: ", beta_6)
```

```
Model coefficients:  [-1.47657400e-10  1.00000000e+00  9.75627389e-14
 6.27429966e-14
 -3.45334740e-15 -2.79829545e-13 -1.15876648e-16 -5.23939616e-14
 1.70411103e-13 -2.59195443e-13 -1.49332369e-13 -3.23608937e-17
 -1.91708477e-21  6.00951580e-15  1.74745461e-11  8.83095680e-14
 6.19702857e-16  1.97735875e-25  6.98209258e-18 -1.12978229e-15
 4.96849675e-18  1.48440288e-14]
```

```
[53]: ### PUNTO 3.
      ### Evalúa con validación cruzada de k-pliegues tu modelo,
      ### calculando los valores de  $R^2$ , MSE y MAE.
```

```

# Residuals
y_pred_6 = predict(X_6, beta_6)
r_6 = y_6 - y_pred_6

# Calculate MSE, MAE and R^2 with the training set
print('MSE: ', mean_squared_error(y_6, y_pred_6))
print("MAE: ", mean_absolute_error(y_6, y_pred_6))
print("R^2: ", r2_score(y_6, y_pred_6))

# Evaluate model with cross validation
n_folds = 5
kf = KFold(n_splits=n_folds, shuffle = True)

mse_cv_6 = []
mae_cv_6 = []
r2_cv_6 = []
for train_index_6, test_index_6 in kf.split(x_6):

    # Training phase
    x_train_6 = x_6[train_index_6, :]
    y_train_6 = y_6[train_index_6]

    beta_cv_6 = fit_model(x_train_6, y_train_6)

    # Test phase
    x_test_6 = x_6[test_index_6, :]
    y_test_6 = y_6[test_index_6]

    y_pred_6 = predict(x_test_6, beta_cv_6)

    # Calculate MSE, MAE and R^2
    mse_i_6 = mean_squared_error(y_test_6, y_pred_6)
    print('mse = ', mse_i_6)
    mse_cv_6.append(mse_i_6)

    mae_i_6 = mean_absolute_error(y_test_6, y_pred_6)
    print('mae = ', mae_i_6)
    mae_cv_6.append(mae_i_6)

    r2_i_6 = r2_score(y_test_6, y_pred_6)
    print('r^2= ', r2_i_6)
    r2_cv_6.append(r2_i_6)

print('MSE:', np.average(mse_cv_6), ' MAE:', np.average(mae_cv_6), ' R^2:', np.
    ↪average(r2_cv_6))

```

MSE: 2.1158358197966936e-22

MAE: 1.1033246414061442e-11


```

R^2: 1.0
mse = 1.8023191537026032e-21
mae = 3.1259746306618165e-11
r^2= 1.0
mse = 2.0960803359932916e-21
mae = 3.340545617902535e-11
r^2= 1.0
mse = 1.9300361220905918e-21
mae = 3.3159780645236783e-11
r^2= 1.0
mse = 1.8644657472353603e-21
mae = 3.166428197084425e-11
r^2= 1.0
mse = 1.3386381541900424e-21
mae = 2.8204205603904313e-11
r^2= 1.0
MSE: 1.806307902642378e-21    MAE: 3.1538694141125766e-11    R^2: 1.0

```

7. Implementa regresión Ridge con descenso de gradiente, y genera el gráfico de Ridge para el conjunto de datos original (sin las variables elevadas al cuadrado).

```
[55]: import numpy.linalg as ln
```

```
[56]: # Rango de valores Lambda
lambdas = np.logspace(-4, 4, 100) # Escala logarítmica de 10^-2 a 10^4

# Vector para almacenamiento de cada lambda
Coeficientes = []

# Ridge regression
def ridge_grad(X, y, beta, lambda_reg):
    n = len(y)
    y_predict = X @ beta
    res = y - y_predict
    grad_mse = -(2 / n) * (X.T @ res)
    grad_ridge = 2 * lambda_reg * beta
    return grad_mse + grad_ridge

# Función para realizar la regresión Ridge utilizando descenso de gradiente
def ridge_regression_gradient_descent(X, y, alpha=0.003, lambda_reg=0.6,
    maxit=10000):
    npredictors = X.shape[1]
    beta = np.random.randn(npredictors) * 0.01
    it = 0

    while (ln.norm(ridge_grad(X, y, beta, lambda_reg)) > 1e-4) and (it < maxit):
        grad = ridge_grad(X, y, beta, lambda_reg)
        grad = np.clip(grad, -1, 1) # Evitamos overflow

```

```

        beta = beta - alpha * grad
        it += 1

    return beta

# Ejecutamos la regression de Ridge
for lambda_reg in lambdas:
    beta_ride = ridge_regression_gradient_descent(X, y, alpha=0.003,
    ↪ lambda_reg=lambda_reg)
    Coeficientes.append(beta_ride)

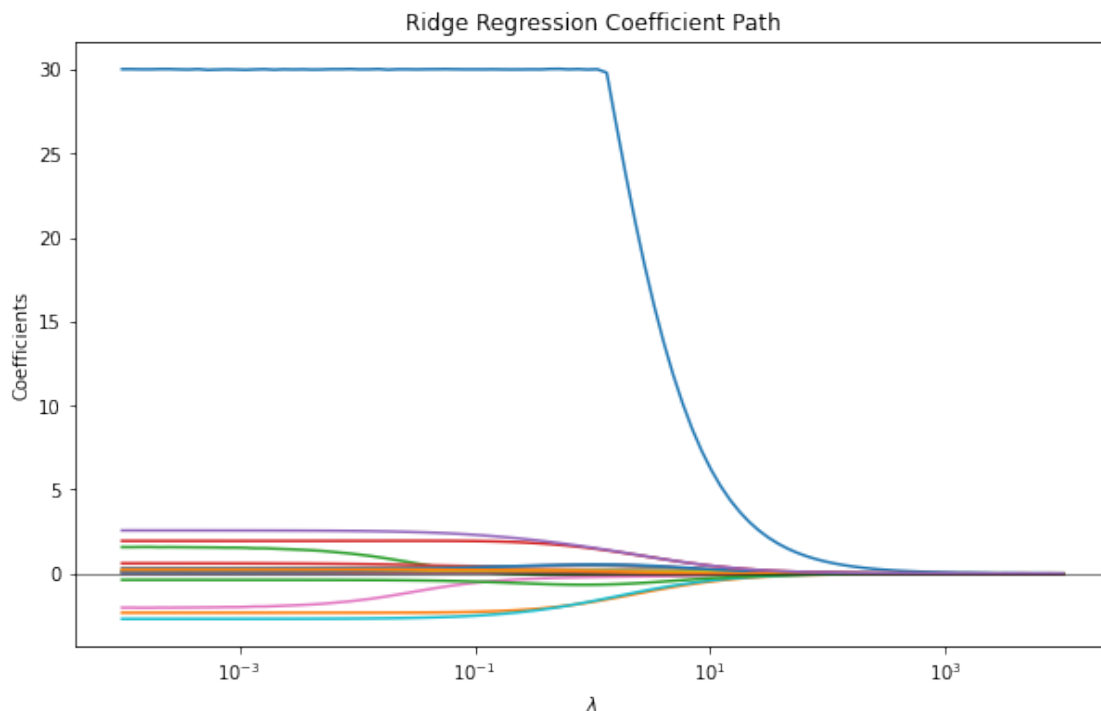
```

```

[57]: Coeficientes = np.array(Coeficientes)

plt.figure(figsize=(10, 6))
for i in range(Coeficientes.shape[1]):
    plt.plot(lambdas, Coeficientes[:, i], label=f'Coef {i+1}')
plt.xscale('log')
plt.xlabel(r'$\lambda$')
plt.ylabel('Coefficients')
plt.title('Ridge Regression Coefficient Path')
plt.axhline(0, color='black', lw=0.5)
plt.show()

```



8. Utiliza una librería para generar el gráfico de Lasso para el conjunto de datos original (sin las variables elevadas al cuadrado). ¿Qué variables son más relevantes para el modelo?

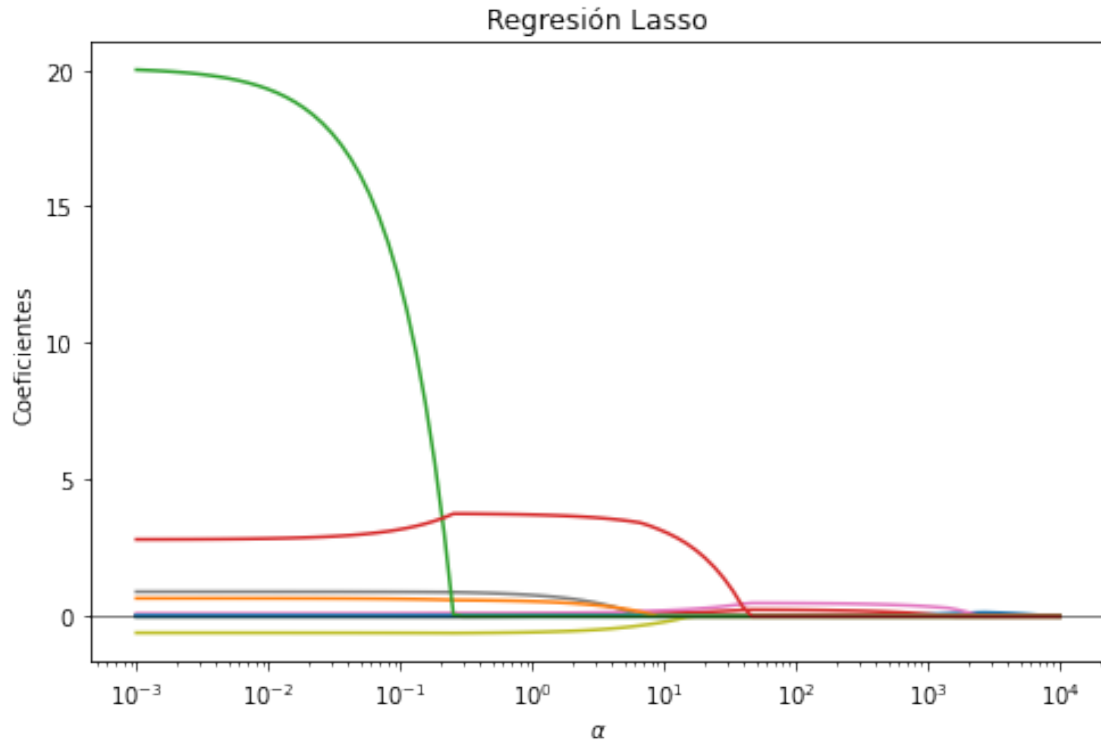
```
[58]: from sklearn.linear_model import Lasso

# Definir un rango de valores para el parámetro de regularización (alpha) en Lasso
alpha_values = np.logspace(-3, 4, 100)
coefficients = []

# Ajustar el modelo Lasso para cada alpha y guardar los coeficientes resultantes
for alpha_value in alpha_values:
    lasso_model = Lasso(alpha=alpha_value, fit_intercept=False, max_iter=10000)
    lasso_model.fit(x, y)
    coefficients.append(lasso_model.coef_)

coefficients = np.array(coefficients)

# Crear la gráfica de los coeficientes en la regresión Lasso
plt.figure(figsize=(8, 5))
for index in range(coefficients.shape[1]):
    plt.plot(alpha_values, coefficients[:, index], label=f'Coef {index + 1}')
plt.xscale('log')
plt.xlabel(r'$\alpha$')
plt.ylabel('Coeficientes')
plt.title('Regresión Lasso')
plt.axhline(0, color='black', linewidth=0.5)
plt.show()
```



9. Viendo los resultados de regresión, desarrolla una conclusión sobre los siguientes puntos:

- (a) ¿Consideras que el modelo de regresión lineal es efectivo para modelar los datos del problema?
¿Por qué?

Si, porque tanto para validación cruzada como Monte Carlo se obtuvo el valor de R^2 muy bueno de 0.82 y los valores de errores de las predicciones, MSE y MAE son pequeños.

- (b) ¿Observas una variabilidad importante en los valores de R^2 , MSE y MAE cuando aplicas validación cruzada?

Realmente no hay una variabilidad importante en las métricas de evaluación. Para MSE es entre 13.422 y 13.3748 sus valores. MAE entre 2.8047 y 2.80189. Como podemos observar es muy insignificante

- (c) ¿Qué modelo es mejor para los datos del problema, el lineal o el cuadrático? ¿Por qué?

Al no haber evaluado con un modelo cuadrático no se puede comparar y definir cuál modelo es mejor para los datos de este problema. Tendríamos que comparar los valores de los errores MSE y MAE, si fuesen menores entonces se pudiese concluir que capta mejor las relaciones entre las variables independientes y dependiente.

- (d) ¿Qué variables son más relevantes para el modelo según Ridge y Lasso?

Para Lasso, mientras alpha reduce los coeficientes tienden a cero haciendo así otras variables más significativas y destacando las que más aportan al modelo. Mientras que para Ridge, conforme

aumenta lambda disminuyen los coeficientes y así resaltan las que aportan también más al modelo.

(e) ¿Encuentras alguna relación interesante entre la variable de respuesta y los predictores?

Que los predictores más significativos son aquellos que tanto en Ridge como en Lasso se ven menos afectados al alterar alpha y lambda, por ende puede representar que tengan mayor relación con la respuesta.

3 EJERCICIO 2

Considere el conjunto de datos de seguimiento telemétrico de la enfermedad de Parkinson (Parkinsons Telemonitoring - UCI Machine Learning Repository), el cual contiene 19 características entre las cuales hay varias derivadas de grabaciones de voz de pacientes con Parkinson. La idea es crear un modelo que prediga a partir de la voz de un paciente la severidad de su enfermedad, la cual es cuantificada con una escala estándar médica llamada UPDRS.

Las características de este conjunto de datos son las siguientes:

- X1 - age
- X2 - test_time
- X3 - Jitter (%)
- X4 - Jitter (Abs)
- X5 - Jitter: RAP
- X6 - Jitter: PPQ5
- X7 - Jitter: DDP
- X8 - Shimmer
- X9- Shimmer (dB)
- X10 - Shimmer: APQ3
- X11 - Shimmer: APQ5
- X12 - Shimmer: APQ11
- X13 - Shimmer: DDA
- X14 - NHR
- X15 - HNR
- X16 - RPDE
- X17 - DFA
- X18 - PPE
- X19 - sex

Como variables dependientes, se tienen **motor_UPDRS** y **total_UPDRS**. Para este ejercicio, se te asignó un conjunto de variables predictoras y una de las variables dependientes de acuerdo a tu matrícula.

*Nota: Las variables con las que vas a trabajar depende del último número de tu matrícula de acuerdo a la siguiente lista: $A01285158 * 8$ - Todas las variables predictoras, menos X4, X8, X12, X16, la variable motor_UPDRS como variable a predecir.*

```
[5]: from ucimlrepo import fetch_ucirepo

# fetch dataset
parkinsons_telemonitoring = fetch_ucirepo(id=189)
```

```
# data (as pandas dataframes)
x = parkinsons_telemonitoring.data.features
y = parkinsons_telemonitoring.data.targets
```

[6]: x

```
[6]:
```

	age	test_time	Jitter(%)	Jitter(Abs)	Jitter:RAP	Jitter:PPQ5	\	
0	72	5.6431	0.00662	0.000034	0.00401	0.00317		
1	72	12.6660	0.00300	0.000017	0.00132	0.00150		
2	72	19.6810	0.00481	0.000025	0.00205	0.00208		
3	72	25.6470	0.00528	0.000027	0.00191	0.00264		
4	72	33.6420	0.00335	0.000020	0.00093	0.00130		
...		
5870	61	142.7900	0.00406	0.000031	0.00167	0.00168		
5871	61	149.8400	0.00297	0.000025	0.00119	0.00147		
5872	61	156.8200	0.00349	0.000025	0.00152	0.00187		
5873	61	163.7300	0.00281	0.000020	0.00128	0.00151		
5874	61	170.7300	0.00282	0.000021	0.00135	0.00166		
...		
	Jitter:DDP	Shimmer	Shimmer(dB)	Shimmer:APQ3	Shimmer:APQ5	\		
0	0.01204	0.02565	0.230	0.01438	0.01309			
1	0.00395	0.02024	0.179	0.00994	0.01072			
2	0.00616	0.01675	0.181	0.00734	0.00844			
3	0.00573	0.02309	0.327	0.01106	0.01265			
4	0.00278	0.01703	0.176	0.00679	0.00929			
...			
5870	0.00500	0.01896	0.160	0.00973	0.01133			
5871	0.00358	0.02315	0.215	0.01052	0.01277			
5872	0.00456	0.02499	0.244	0.01371	0.01456			
5873	0.00383	0.01484	0.131	0.00693	0.00870			
5874	0.00406	0.01907	0.171	0.00946	0.01154			
...			
	Shimmer:APQ11	Shimmer:DDA	NHR	HNR	RPDE	DFA	PPE	\
0	0.01662	0.04314	0.014290	21.640	0.41888	0.54842	0.16006	
1	0.01689	0.02982	0.011112	27.183	0.43493	0.56477	0.10810	
2	0.01458	0.02202	0.020220	23.047	0.46222	0.54405	0.21014	
3	0.01963	0.03317	0.027837	24.445	0.48730	0.57794	0.33277	
4	0.01819	0.02036	0.011625	26.126	0.47188	0.56122	0.19361	
...	
5870	0.01549	0.02920	0.025137	22.369	0.64215	0.55314	0.21367	
5871	0.01904	0.03157	0.011927	22.886	0.52598	0.56518	0.12621	
5872	0.01877	0.04112	0.017701	25.065	0.47792	0.57888	0.14157	
5873	0.01307	0.02078	0.007984	24.422	0.56865	0.56327	0.14204	
5874	0.01470	0.02839	0.008172	23.259	0.58608	0.57077	0.15336	

sex

```

0      0
1      0
2      0
3      0
4      0
...
5870   0
5871   0
5872   0
5873   0
5874   0

```

[5875 rows x 19 columns]

```
[7]: x = x.drop(['Jitter(Abs)', 'Shimmer', 'Shimmer:APQ11', 'RPDE'], axis=1)
x
```

```
[7]:
```

	age	test_time	Jitter(%)	Jitter:RAP	Jitter:PPQ5	Jitter:DDP	\
0	72	5.6431	0.00662	0.00401	0.00317	0.01204	
1	72	12.6660	0.00300	0.00132	0.00150	0.00395	
2	72	19.6810	0.00481	0.00205	0.00208	0.00616	
3	72	25.6470	0.00528	0.00191	0.00264	0.00573	
4	72	33.6420	0.00335	0.00093	0.00130	0.00278	
...	
5870	61	142.7900	0.00406	0.00167	0.00168	0.00500	
5871	61	149.8400	0.00297	0.00119	0.00147	0.00358	
5872	61	156.8200	0.00349	0.00152	0.00187	0.00456	
5873	61	163.7300	0.00281	0.00128	0.00151	0.00383	
5874	61	170.7300	0.00282	0.00135	0.00166	0.00406	
...	
	Shimmer(dB)	Shimmer:APQ3	Shimmer:APQ5	Shimmer:DDA	NHR	HNR	\
0	0.230	0.01438	0.01309	0.04314	0.014290	21.640	
1	0.179	0.00994	0.01072	0.02982	0.011112	27.183	
2	0.181	0.00734	0.00844	0.02202	0.020220	23.047	
3	0.327	0.01106	0.01265	0.03317	0.027837	24.445	
4	0.176	0.00679	0.00929	0.02036	0.011625	26.126	
...	
5870	0.160	0.00973	0.01133	0.02920	0.025137	22.369	
5871	0.215	0.01052	0.01277	0.03157	0.011927	22.886	
5872	0.244	0.01371	0.01456	0.04112	0.017701	25.065	
5873	0.131	0.00693	0.00870	0.02078	0.007984	24.422	
5874	0.171	0.00946	0.01154	0.02839	0.008172	23.259	
...	
	DFA	PPE	sex				
0	0.54842	0.16006	0				
1	0.56477	0.10810	0				
2	0.54405	0.21014	0				

3	0.57794	0.33277	0
4	0.56122	0.19361	0
...
5870	0.55314	0.21367	0
5871	0.56518	0.12621	0
5872	0.57888	0.14157	0
5873	0.56327	0.14204	0
5874	0.57077	0.15336	0

[5875 rows x 15 columns]

```
[8]: y
```

```
[8]:      motor_UPDRS  total_UPDRS
0          28.199        34.398
1          28.447        34.894
2          28.695        35.389
3          28.905        35.810
4          29.187        36.375
...
5870        22.485        33.485
5871        21.988        32.988
5872        21.495        32.495
5873        21.007        32.007
5874        20.513        31.513
```

[5875 rows x 2 columns]

```
[9]: y = y.drop(['total_UPDRS'],axis=1)
y
```

```
[9]:      motor_UPDRS
0          28.199
1          28.447
2          28.695
3          28.905
4          29.187
...
5870        22.485
5871        21.988
5872        21.495
5873        21.007
5874        20.513
```

[5875 rows x 1 columns]


```
[10]: x = x.to_numpy()
      y = y.to_numpy()
```

Para este conjunto de datos y las variables que se te asignaron:

1. Evalúa con validación cruzada un modelo de regresión lineal para las variables asignadas según tu matrícula utilizando alguna librería o framework.

```
[11]: import numpy as np
      import matplotlib.pyplot as plt
      from sklearn import datasets
      from sklearn import linear_model
      from sklearn.model_selection import KFold
      from sklearn.feature_selection import SelectKBest, r_regression
      from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
      import warnings
      warnings.filterwarnings("ignore")
```

```
[12]: ##### Train linear regression model
      regr = linear_model.LinearRegression()
      regr.fit(x, y)
      print("Coeficientes del modelo: ", regr.coef_)
      print("Intercepto del modelo : ", regr.intercept_)

      y_pred = regr.predict(x)
      print('MSE: ', mean_squared_error(y, y_pred))
      print("MAE: ", mean_absolute_error(y, y_pred))
      print("R^2: ", r2_score(y, y_pred))

      # k-fold cross-validation
      n_folds = 5
      kf = KFold(n_splits=n_folds, shuffle = True)

      mse_cv = []
      mae_cv = []
      r2_cv = []
      for train_index, test_index in kf.split(x):

          # Training phase
          x_train = x[train_index, :]
          y_train = y[train_index]

          fselection_cv = SelectKBest(r_regression, k=5)
          fselection_cv.fit(x_train, y_train)
          x_train = fselection_cv.transform(x_train)

          regr_cv = linear_model.LinearRegression()
          regr_cv.fit(x_train, y_train)
```

```

# Test phase
x_test = fselection_cv.transform(x[test_index, :])
y_test = y[test_index]
y_pred = regr_cv.predict(x_test)

mse_i = mean_squared_error(y_test, y_pred)
mse_cv.append(mse_i)

mae_i = mean_absolute_error(y_test, y_pred)
mae_cv.append(mae_i)

r2_i = r2_score(y_test, y_pred)
r2_cv.append(r2_i)

print('MSE:', np.average(mse_cv), ' MAE:', np.average(mae_cv), ' R^2:', np.
↪average(r2_cv))

```

```

Coeficientes del modelo: [[ 1.96998903e-01  1.10274984e-02 -8.63687865e+01
-4.19622234e+04
 2.46148191e+01  1.40778877e+04  9.05626170e+00  2.99504084e+03
 2.92448294e+01 -1.07243681e+03 -2.52432651e+01 -4.20562133e-01
-2.60895353e+01  1.20988987e+01 -3.87986109e-01]]

```

```
Intercepto del modelo : [31.98931322]
```

```
MSE: 56.78867312476895
```

```
MAE: 6.385296161246183
```

```
R^2: 0.14052881169564913
```

```
MSE: 59.91026483227439 MAE: 6.643827561426397 R^2: 0.0927936385266388
```

2. Encuentra el número óptimo de predictores para el modelo utilizando el método filter y validación cruzada. Una vez que tengas el número óptimo, muestra las características seleccionadas.

```

[68]: #-----
# Find optimal number of features using cross-validation
# FILTER
#-----
print("----- Optimal selection of number of features -----")
print("----- FILTER SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

```

```

mse_cv = []
mae_cv = []
r2_cv = []

kf = KFold(n_splits=5, shuffle = True)

for train_index, test_index in kf.split(x):

    # Training phase
    x_train = x[train_index, :]
    y_train = y[train_index]

    fselection_cv = SelectKBest(r_regression, k = n_feat)
    fselection_cv.fit(x_train, y_train)
    x_train = fselection_cv.transform(x_train)

    regr_cv = linear_model.LinearRegression()
    regr_cv.fit(x_train, y_train)

    # Test phase
    x_test = fselection_cv.transform(x[test_index, :])
    y_test = y[test_index]
    y_pred = regr_cv.predict(x_test)

    mse_i = mean_squared_error(y_test, y_pred)
    mse_cv.append(mse_i)

    mae_i = mean_absolute_error(y_test, y_pred)
    mae_cv.append(mae_i)

    r2_i = r2_score(y_test, y_pred)
    r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]

```

```

print("Optimal number of features: ", opt_features)

----- Optimal selection of number of features -----
----- FILTER SELECTION -----
---- n features = 1
MSE: 61.174946509667485    MAE: 6.715128948821858    R^2: 0.07326818239968942
---- n features = 2
MSE: 60.050144811211354    MAE: 6.6503617166118065    R^2: 0.09004053181187661
---- n features = 3
MSE: 60.09043872349211    MAE: 6.652426971723817    R^2: 0.0896033777707405
---- n features = 4
MSE: 60.03493223670095    MAE: 6.652018277708114    R^2: 0.09072139443298899
---- n features = 5
MSE: 59.87388903525086    MAE: 6.642376490676111    R^2: 0.09359996450352528
---- n features = 6
MSE: 59.837694382954794    MAE: 6.6386310973545335    R^2: 0.09400773710259451
---- n features = 7
MSE: 59.864973859892544    MAE: 6.639217048533408    R^2: 0.09308651882097033
---- n features = 8
MSE: 59.820675021971645    MAE: 6.648832257824685    R^2: 0.0943911549567324
---- n features = 9
MSE: 59.798716690428115    MAE: 6.645503361964613    R^2: 0.09476039115679022
Optimal number of features:  9

```

```

[69]: fig, axs = plt.subplots(1, 3, tight_layout=True)
      axs[0].plot(n_feats, mse_nfeat)
      axs[0].set_xlabel("k")
      axs[0].set_ylabel("MSE")

      axs[1].plot(n_feats, mae_nfeat)
      axs[1].set_xlabel("k")
      axs[1].set_ylabel("MAE")

      axs[2].plot(n_feats, r2_nfeat)
      axs[2].set_xlabel("k")
      axs[2].set_ylabel("r^2")

      plt.show()

      # Fit model with optimal number of features
      regr = linear_model.LinearRegression()
      fselection = SelectKBest(r_regression, k = opt_features)
      fselection.fit(x, y)

      print("Selected features: ", fselection.get_feature_names_out())

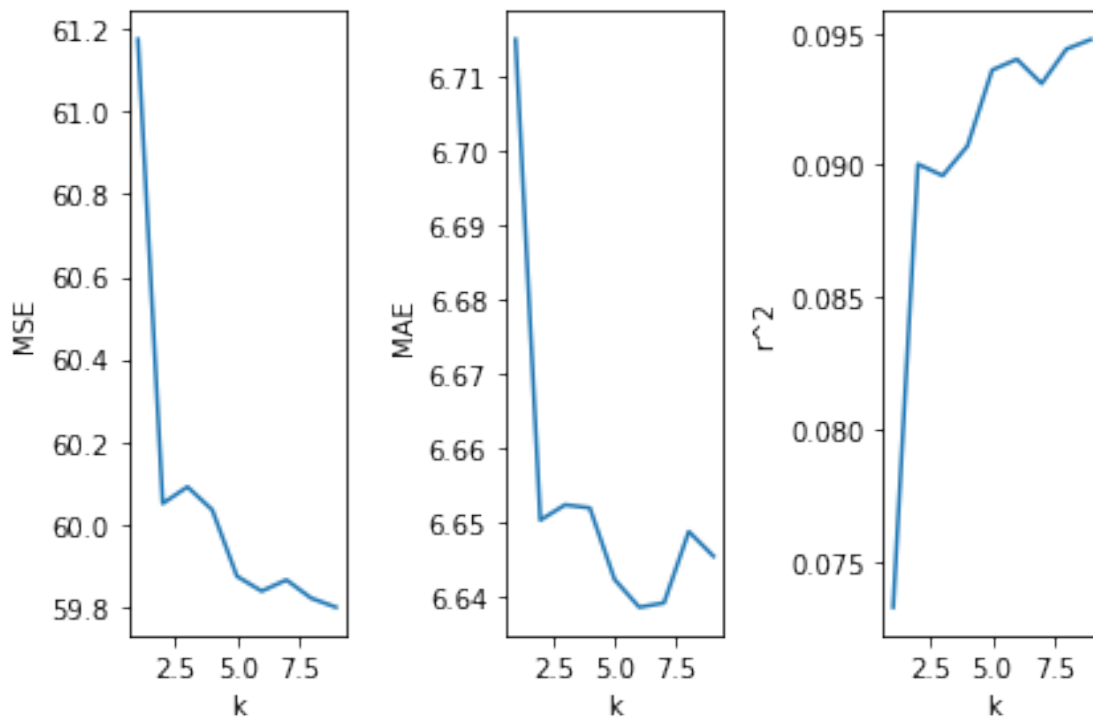
      x_transformed = fselection.transform(x)

```

```

regr.fit(x_transformed, y)
print("Model coefficients: ", regr.coef_)
print("Model intercept: ", regr.intercept_)

```



```

Selected features: ['x0' 'x2' 'x4' 'x6' 'x7' 'x8' 'x9' 'x10' 'x13']
Model coefficients: [[ 2.32442119e-01 -8.19405000e+01 -1.24217423e+01
 1.11732287e+01
 -3.78708493e+02 -3.74766014e+01  7.94157967e+01  2.94833521e+00
 1.25224574e+01]]
Model intercept: [3.62351321]

```

3. Repite el paso anterior pero con selección de características secuencial (Wrapper). Reporta los predictores óptimos encontrados por el método.

```

[70]: from sklearn.feature_selection import SequentialFeatureSelector

#####
# Find optimal number of features using cross-validation
#####
print("----- Optimal selection of number of features -----")
print("----- WRAPPER SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

```

```

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)

    for train_index, test_index in kf.split(x):

        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]

        regr_cv = linear_model.LinearRegression()

        fselection_cv = SequentialFeatureSelector(regr_cv,
↪n_features_to_select=n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)

        regr_cv.fit(x_train, y_train)

        # Test phase
        x_test = fselection_cv.transform(x[test_index, :])
        y_test = y[test_index]
        y_pred = regr_cv.predict(x_test)

        mse_i = mean_squared_error(y_test, y_pred)
        mse_cv.append(mse_i)

        mae_i = mean_absolute_error(y_test, y_pred)
        mae_cv.append(mae_i)

        r2_i = r2_score(y_test, y_pred)
        r2_cv.append(r2_i)

    mse = np.average(mse_cv)
    mse_nfeat.append(mse)

    mae = np.average(mae_cv)
    mae_nfeat.append(mae)

```

```

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

----- Optimal selection of number of features -----
----- WRAPPER SELECTION -----
---- n features = 1
MSE: 61.16136215066884    MAE: 6.714880851525959    R^2: 0.07378184291508232
---- n features = 2
MSE: 60.06734185765474    MAE: 6.65047366521228    R^2: 0.08993583048279433
---- n features = 3
MSE: 60.01152135970888    MAE: 6.64654947120294    R^2: 0.09147629703427604
---- n features = 4
MSE: 60.0830552639873    MAE: 6.65003142190194    R^2: 0.08937716771953294
---- n features = 5
MSE: 60.046736499097335    MAE: 6.646313875108554    R^2: 0.09082264581586255
---- n features = 6
MSE: 60.15524211708055    MAE: 6.653076426991136    R^2: 0.08826122368458311
---- n features = 7
MSE: 60.09135294589309    MAE: 6.644698565052657    R^2: 0.0900030423903243
---- n features = 8
MSE: 59.99969842234283    MAE: 6.6751435506053625    R^2: 0.09143128410757048
---- n features = 9
MSE: 59.948591021894586    MAE: 6.675300380942545    R^2: 0.09154425150660565
Optimal number of features:  9

```

```

[71]: fig, axs = plt.subplots(1, 3, tight_layout=True)
      axs[0].plot(n_feats, mse_nfeat)
      axs[0].set_xlabel("features")
      axs[0].set_ylabel("MSE")

      axs[1].plot(n_feats, mae_nfeat)
      axs[1].set_xlabel("features")
      axs[1].set_ylabel("MAE")

      axs[2].plot(n_feats, r2_nfeat)
      axs[2].set_xlabel("features")
      axs[2].set_ylabel("r^2")

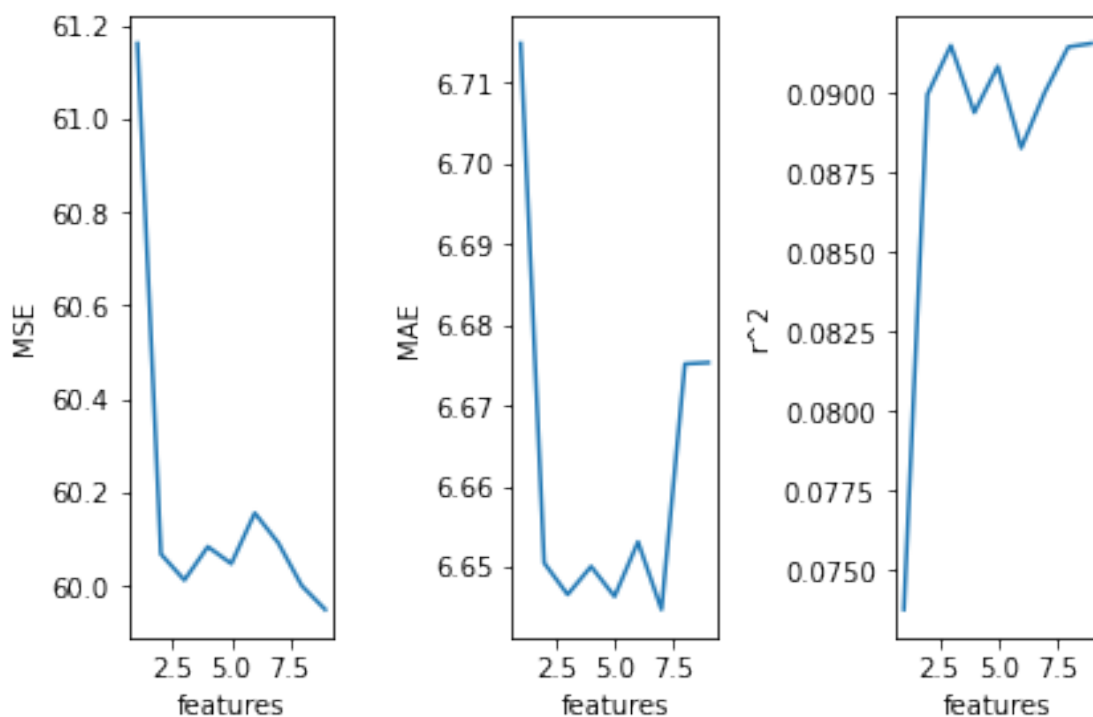
      plt.show()

```

```
# Fit model with optimal number of features
regr = linear_model.LinearRegression()
fselection = SequentialFeatureSelector(regr, n_features_to_select = 10,
    step_forward=1, step_backward=1)
fselection.fit(x, y)

print("Selected features: ", fselection.get_feature_names_out())

x_transformed = fselection.transform(x)
regr.fit(x_transformed, y)
print("Model coefficients: ", regr.coef_)
print("Model intercept: ", regr.intercept_)
```



```
Selected features: ['x0' 'x1' 'x2' 'x3' 'x4' 'x5' 'x6' 'x10' 'x13']
Model coefficients: [[ 2.34916565e-01  9.47142993e-03 -6.85631774e+01
 -4.01117996e+04
 -3.25001421e+01  1.33650251e+04 -9.95783776e-02  6.28827924e+00
  1.38316113e+01]]
Model intercept: [2.56137755]
```

4. Haz el mismo proceso del paso 2, pero ahora con el método de selección de características recursivo. Reporta los predictores óptimos encontrados por el método.


```
[72]: #-----
# Recursive feature selection
#-----

from sklearn.feature_selection import RFE

#####
# Find optimal number of features using cross-validation
#####
print("----- Optimal selection of number of features -----")
print("----- RECURSIVE SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)

    for train_index, test_index in kf.split(x):

        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]

        regr_cv = linear_model.LinearRegression()

        fselection_cv = RFE(regr_cv, n_features_to_select=n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)

        regr_cv.fit(x_train, y_train)

        # Test phase
        x_test = fselection_cv.transform(x[test_index, :])
        y_test = y[test_index]
        y_pred = regr_cv.predict(x_test)

        mse_i = mean_squared_error(y_test, y_pred)
        mse_cv.append(mse_i)
```

```

    mae_i = mean_absolute_error(y_test, y_pred)
    mae_cv.append(mae_i)

    r2_i = r2_score(y_test, y_pred)
    r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

----- Optimal selection of number of features -----
----- RECURSIVE SELECTION -----
---- n features = 1
MSE: 65.76510884382762    MAE: 6.950878760185704    R^2: 0.0037645558470920993
---- n features = 2
MSE: 65.69962489210096    MAE: 6.940458446126646    R^2: 0.00553183234179524
---- n features = 3
MSE: 65.59125984812411    MAE: 6.927069140530925    R^2: 0.007247103834636603
---- n features = 4
MSE: 65.6818558321831    MAE: 6.931594033247256    R^2: 0.005021547819924565
---- n features = 5
MSE: 65.35601834809556    MAE: 6.931107299592901    R^2: 0.010672286929717023
---- n features = 6
MSE: 64.7466936803547    MAE: 6.878076775382262    R^2: 0.01987334358959658
---- n features = 7
MSE: 65.12631580148609    MAE: 6.90778077505212    R^2: 0.01264604608304869
---- n features = 8
MSE: 63.13304761610591    MAE: 6.774519896707392    R^2: 0.04286883153991052
---- n features = 9
MSE: 61.44210076592229    MAE: 6.659175637217873    R^2: 0.06877877478590733
Optimal number of features: 9

```

```

[73]: fig, axs = plt.subplots(1, 3, tight_layout=True)
      axs[0].plot(n_feats, mse_nfeat)

```

```

axs[0].set_xlabel("features")
axs[0].set_ylabel("MSE")

axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("features")
axs[1].set_ylabel("MAE")

axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("features")
axs[2].set_ylabel("r^2")

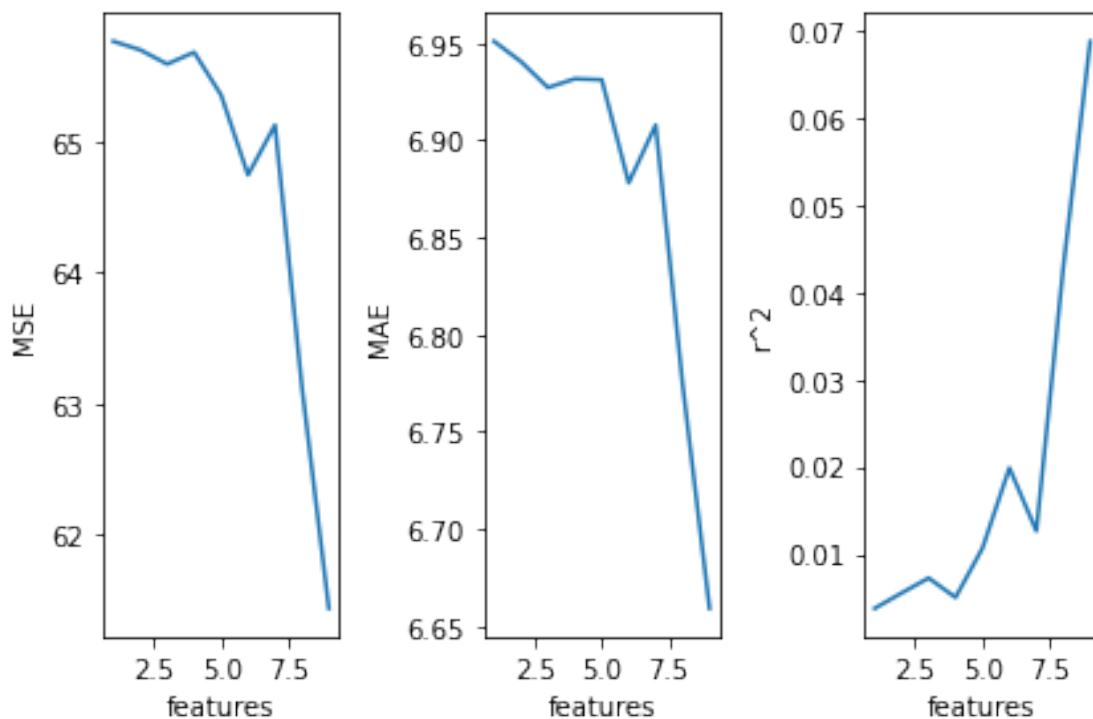
plt.show()

# Fit model with optimal number of features
regr = linear_model.LinearRegression()
fselection = RFE(regr, n_features_to_select = opt_features)
fselection.fit(x, y)

print("Selected features: ", fselection.get_feature_names_out())

x_transformed = fselection.transform(x)
regr.fit(x_transformed, y)
print("Model coefficients: ", regr.coef_)
print("Model intercept: ", regr.intercept_)

```



```

Selected features: ['x2' 'x3' 'x4' 'x5' 'x7' 'x8' 'x9' 'x12' 'x13']
Model coefficients: [[-4.21734714e+01 -5.27340468e+04 -3.02359779e+02
1.76322592e+04
-5.87995771e+03 6.90510348e+01 1.93321384e+03 -2.56332454e+01
2.78140914e+01]]
Model intercept: [32.68248611]

```

5. Repita los pasos anteriores, pero utilizando un modelo de regresión no lineal como K-vecinos más cercanos.

```
[74]: from sklearn.neighbors import KNeighborsRegressor
```

METODO FILTER

```
[75]: #-----
# Find optimal number of features using cross-validation
# FILTER
#-----
print("----- Optimal selection of number of features -----")
print("----- FILTER SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)

    for train_index, test_index in kf.split(x):

        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]

        fselection_cv = SelectKBest(r_regression, k = n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)

        regr_cv = KNeighborsRegressor(n_neighbors=5) # LinearRegression --> KNN
        regr_cv.fit(x_train, y_train)
```

```

# Test phase
x_test = fselection_cv.transform(x[test_index, :])
y_test = y[test_index]
y_pred = regr_cv.predict(x_test)

mse_i = mean_squared_error(y_test, y_pred)
mse_cv.append(mse_i)

mae_i = mean_absolute_error(y_test, y_pred)
mae_cv.append(mae_i)

r2_i = r2_score(y_test, y_pred)
r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

```

----- Optimal selection of number of features -----

----- FILTER SELECTION -----

---- n features = 1	MSE: 48.29611253522526	MAE: 4.750310498723404	R ² : 0.268088081187167
---- n features = 2	MSE: 24.915874830676696	MAE: 3.54140331574468	R ² : 0.6227272390390443
---- n features = 3	MSE: 22.289202849171268	MAE: 3.272341606808511	R ² : 0.661506896435775
---- n features = 4	MSE: 22.417087654280785	MAE: 3.2881933923404256	R ² : 0.660545339833124
---- n features = 5	MSE: 21.459499767300223	MAE: 3.233451901276596	R ² : 0.6748929172250986
---- n features = 6	MSE: 22.220282366385156	MAE: 3.263428738723404	R ² : 0.6629976297391164
---- n features = 7	MSE: 21.211810581183794	MAE: 3.209076333617021	R ² : 0.6788600977660589
---- n features = 8	MSE: 21.86354376128967	MAE: 3.2383500187234042	R ² : 0.6689965918979915

```
---- n features = 9
MSE: 16.011292366530043    MAE: 2.5931766604255317    R^2: 0.7563930223375073
Optimal number of features: 9
```

```
[76]: fig, axs = plt.subplots(1, 3, tight_layout=True)
      axs[0].plot(n_feats, mse_nfeat)
      axs[0].set_xlabel("k")
      axs[0].set_ylabel("MSE")

      axs[1].plot(n_feats, mae_nfeat)
      axs[1].set_xlabel("k")
      axs[1].set_ylabel("MAE")

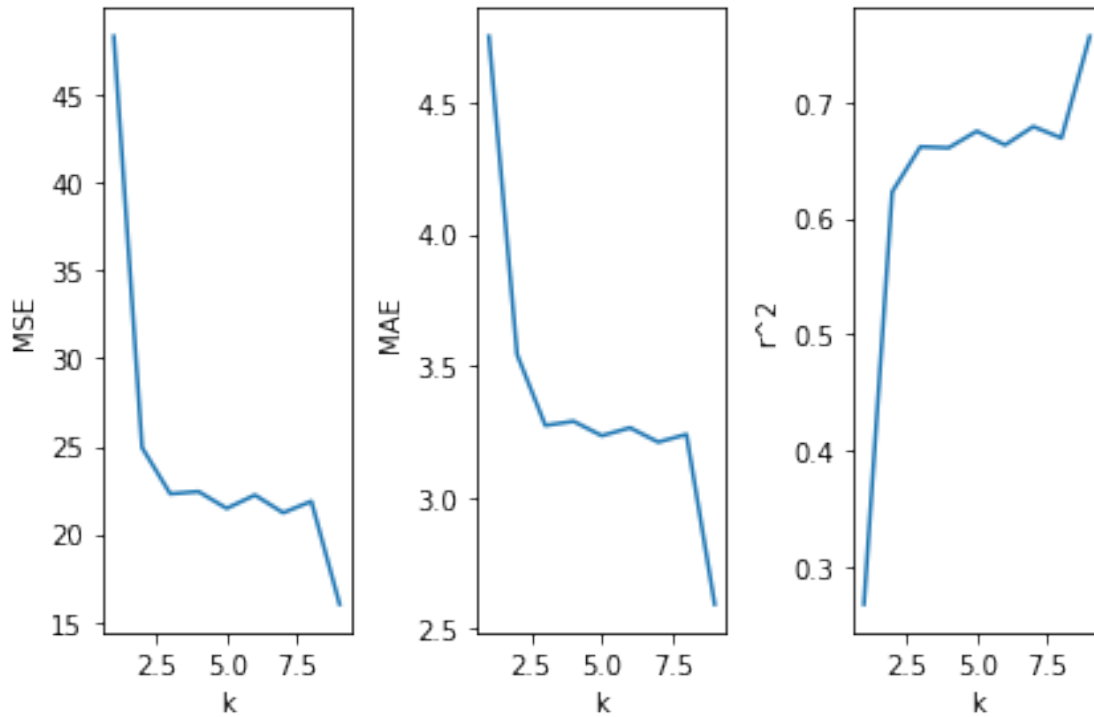
      axs[2].plot(n_feats, r2_nfeat)
      axs[2].set_xlabel("k")
      axs[2].set_ylabel("r^2")

      plt.show()

      # Fit model with optimal number of features
      regr = KNeighborsRegressor() # linear_model.LinearRegression() -> KNN
      fselection = SelectKBest(r_regression, k = opt_features)
      fselection.fit(x, y)

      print("Selected features: ", fselection.get_feature_names_out())

      x_transformed = fselection.transform(x)
      regr.fit(x_transformed, y)
```



Selected features: ['x0' 'x2' 'x4' 'x6' 'x7' 'x8' 'x9' 'x10' 'x13']

[76]: KNeighborsRegressor()

METODO WRAPPER

```
[77]: #####
# Find optimal number of features using cross-validation
#####
print("----- Optimal selection of number of features -----")
print("----- WRAPPER SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)
```

```

for train_index, test_index in kf.split(x):

    # Training phase
    x_train = x[train_index, :]
    y_train = y[train_index]

    regr_cv = KNeighborsRegressor(n_neighbors=5) # LINEAR REGRESSION -> KNN

    fselection_cv = SequentialFeatureSelector(regr_cv, u
↪n_features_to_select=n_feat)
    fselection_cv.fit(x_train, y_train)
    x_train = fselection_cv.transform(x_train)

    regr_cv.fit(x_train, y_train)

    # Test phase
    x_test = fselection_cv.transform(x[test_index, :])
    y_test = y[test_index]
    y_pred = regr_cv.predict(x_test)

    mse_i = mean_squared_error(y_test, y_pred)
    mse_cv.append(mse_i)

    mae_i = mean_absolute_error(y_test, y_pred)
    mae_cv.append(mae_i)

    r2_i = r2_score(y_test, y_pred)
    r2_cv.append(r2_i)

    mse = np.average(mse_cv)
    mse_nfeat.append(mse)

    mae = np.average(mae_cv)
    mae_nfeat.append(mae)

    r2 = np.average(r2_cv)
    r2_nfeat.append(r2)

    print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

```

```

----- Optimal selection of number of features -----
----- WRAPPER SELECTION -----

```



```

---- n features = 1
MSE: 60.246674442795644    MAE: 5.801030852765957    R^2: 0.08759882712984568
---- n features = 2
MSE: 33.34320816772011    MAE: 4.138727281702128    R^2: 0.49615911128646556
---- n features = 3
MSE: 23.831305453509174    MAE: 3.2459111080851066    R^2: 0.6454004422198045
---- n features = 4
MSE: 10.691616195588425    MAE: 2.2728416306382977    R^2: 0.8381773336940848
---- n features = 5
MSE: 10.282840889221651    MAE: 2.2103280102127663    R^2: 0.8445147456388995
---- n features = 6
MSE: 22.72019364512606    MAE: 3.0970031148936172    R^2: 0.6537347385359589
---- n features = 7
MSE: 22.175839089903864    MAE: 3.0476611029787235    R^2: 0.6654285480554479
---- n features = 8
MSE: 10.404622455322961    MAE: 2.2247268221276597    R^2: 0.8425063634020491
---- n features = 9
MSE: 21.23577970423278    MAE: 3.0061730621276594    R^2: 0.6766973619114662
Optimal number of features: 5

```

```

[78]: fig, axs = plt.subplots(1, 3, tight_layout=True)
      axs[0].plot(n_feats, mse_nfeat)
      axs[0].set_xlabel("k")
      axs[0].set_ylabel("MSE")

      axs[1].plot(n_feats, mae_nfeat)
      axs[1].set_xlabel("k")
      axs[1].set_ylabel("MAE")

      axs[2].plot(n_feats, r2_nfeat)
      axs[2].set_xlabel("k")
      axs[2].set_ylabel("r^2")

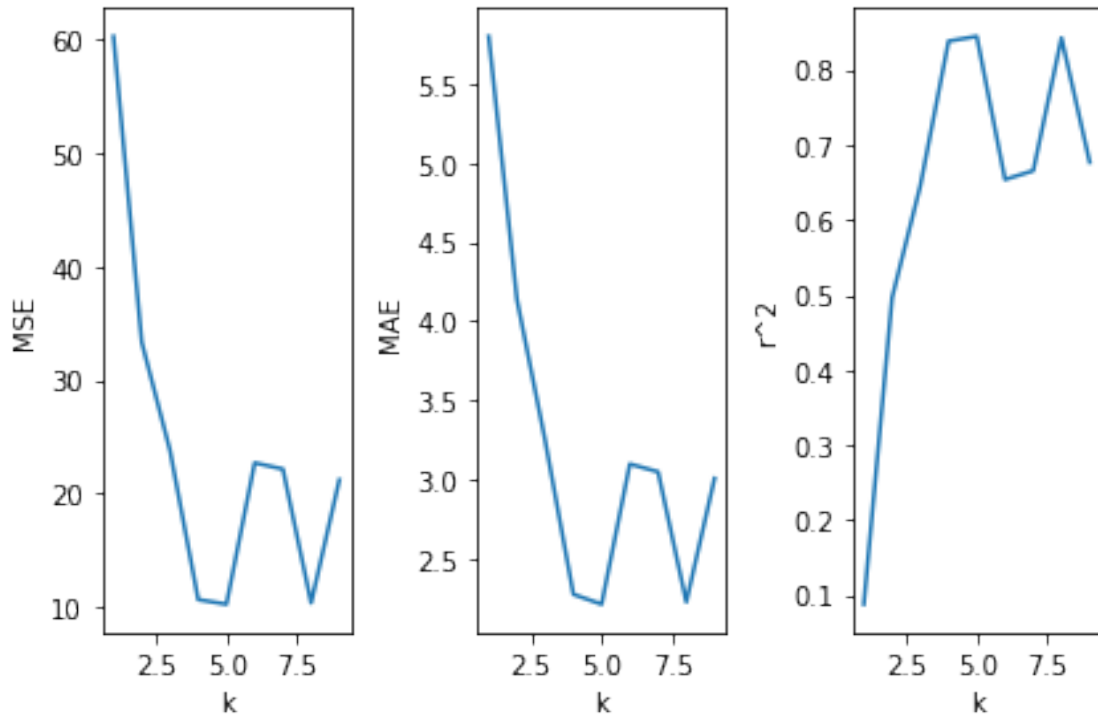
      plt.show()

      # Fit model with optimal number of features using KNN
      regr = KNeighborsRegressor(n_neighbors=5) # LINEAR R --> KNN
      fselection = SelectKBest(r_regression, k=opt_features)
      fselection.fit(x, y)

      print("Selected features: ", fselection.get_feature_names_out())

      x_transformed = fselection.transform(x)
      regr.fit(x_transformed, y)

```



Selected features: ['x0' 'x2' 'x6' 'x8' 'x13']

[78]: KNeighborsRegressor()

6. Busca al menos otros 4 modelos de regresión no lineal, y lleva a cabo los pasos del 1 al 5. Reminder PASOS:

1. Evalúa con validación cruzada un modelo de regresión lineal para las variables asignadas según tu matrícula utilizando alguna librería o framework.
2. Encuentra el número óptimo de predictores para el modelo utilizando el método filter y validación cruzada. Una vez que tengas el número óptimo, muestra las características seleccionadas.
3. Repite el paso anterior pero con selección de características secuencial (Wrapper). Reporta los predictores óptimos encontrados por el método.
4. Haz el mismo proceso del paso 2, pero ahora con el método de selección de características recursivo. Reporta los predictores óptimos encontrados por el método.
5. Repita los pasos anteriores, pero utilizando un modelo de regresión no lineal como K-vecinos más cercanos.

GRACE: Eliminamos el Paso 1 y 5, para este Paso 6 dado que:

- El paso 1 pide una reregresión lineal y queremos ahora no lineales
- El paso 5 ya hace una regresión no lineal (kNN), no vale la pena repetirlo

Nos quedamos con el 2 -> A , 3 -> B, 4 -> C

MODELO 1. GRADIENT BOOSTING REGRESSOR

A. Encuentra el número óptimo de predictores para el modelo utilizando el método filter y validación cruzada. Una vez que tengas el número óptimo, muestra las características seleccionadas.

```
[82]: from sklearn.ensemble import GradientBoostingRegressor

#-----
# Find optimal number of features using cross-validation
# FILTER
#-----

print("----- Optimal selection of number of features -----")
print("----- FILTER SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)

    for train_index, test_index in kf.split(x):

        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]

        fselection_cv = SelectKBest(r_regression, k = n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)

        regr_cv = GradientBoostingRegressor()
        regr_cv.fit(x_train, y_train)

        # Test phase
        x_test = fselection_cv.transform(x[test_index, :])
        y_test = y[test_index]
        y_pred = regr_cv.predict(x_test)

    mse_i = mean_squared_error(y_test, y_pred)
```

```

mse_cv.append(mse_i)

mae_i = mean_absolute_error(y_test, y_pred)
mae_cv.append(mae_i)

r2_i = r2_score(y_test, y_pred)
r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

```

----- Optimal selection of number of features -----
----- FILTER SELECTION -----

n features	MSE	MAE	R ²
1	24.764564791618273	3.9853370181399392	0.6249963857467588
2	24.153020737164717	3.932470671673518	0.6340792201016352
3	23.955042437879914	3.938692961316274	0.6370795791219027
4	23.12526950522358	3.857064214204351	0.6499117522740445
5	23.618955261547747	3.9229283667432613	0.6421920077201734
6	23.431237168962777	3.9089362617472814	0.6453109231790506
7	23.505640786477212	3.922387810270041	0.6439411909752183
8	23.10606284298398	3.8636599184313676	0.650314575655824
9	22.86365673903294	3.8462543591529554	0.6536126046733083

Optimal number of features: 9

```
[83]: fig, axs = plt.subplots(1, 3, tight_layout=True)
axs[0].plot(n_feats, mse_nfeat)
axs[0].set_xlabel("k")
axs[0].set_ylabel("MSE")

axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("k")
axs[1].set_ylabel("MAE")

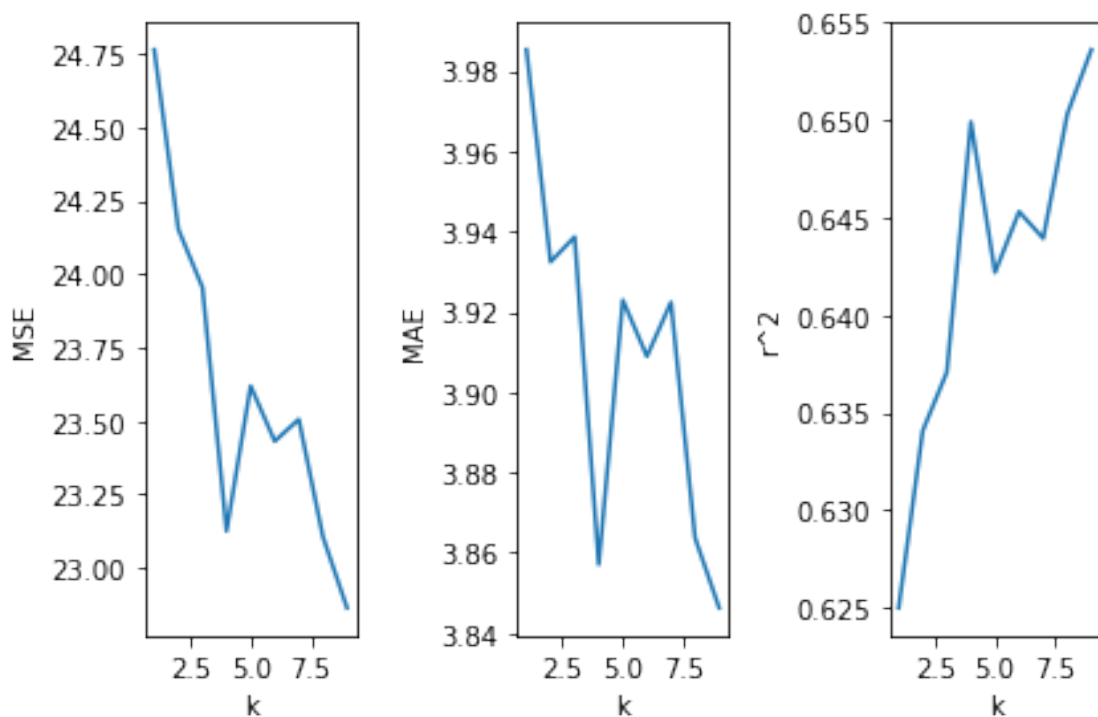
axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("k")
axs[2].set_ylabel("r2")

plt.show()

# Fit model with optimal number of features
regr = GradientBoostingRegressor()
fselection = SelectKBest(r_regression, k = opt_features)
fselection.fit(x, y)

print("Selected features: ", fselection.get_feature_names_out())

x_transformed = fselection.transform(x)
regr.fit(x_transformed, y)
print("Feature importances:", regr.feature_importances_)
```



```

Selected features: ['x0' 'x2' 'x4' 'x6' 'x7' 'x8' 'x9' 'x10' 'x13']
Feature importances: [0.929131  0.01031222 0.00693149 0.00310831 0.00251368
0.01276572
0.0055333  0.01682146 0.01288283]

```

B. Repite el paso anterior pero con selección de características secuencial (Wrapper). Reporta los predictores óptimos encontrados por el método.

```

[84]: from sklearn.ensemble import GradientBoostingRegressor

#####
# Find optimal number of features using cross-validation
#####
print("----- Optimal selection of number of features -----")
print("----- WRAPPER SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('----- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)

    for train_index, test_index in kf.split(x):

        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]

        regr_cv = GradientBoostingRegressor()

        fselection_cv = SequentialFeatureSelector(regr_cv,
↪n_features_to_select=n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)

        regr_cv.fit(x_train, y_train)

```

```

# Test phase
x_test = fselection_cv.transform(x[test_index, :])
y_test = y[test_index]
y_pred = regr_cv.predict(x_test)

mse_i = mean_squared_error(y_test, y_pred)
mse_cv.append(mse_i)

mae_i = mean_absolute_error(y_test, y_pred)
mae_cv.append(mae_i)

r2_i = r2_score(y_test, y_pred)
r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

```

----- Optimal selection of number of features -----

----- WRAPPER SELECTION -----

---- n features = 1	MSE: 65.07859163983865	MAE: 6.878019949102869	R^2: 0.014122444789345124
---- n features = 2	MSE: 47.70782216933044	MAE: 5.569030477429966	R^2: 0.285371547966795
---- n features = 3	MSE: 60.206879939076615	MAE: 6.551982361778473	R^2: 0.08842872183264247
---- n features = 4	MSE: 51.38918087513417	MAE: 5.872266154574374	R^2: 0.21521201870386192
---- n features = 5	MSE: 52.08286554891821	MAE: 5.9105178477405556	R^2: 0.21720347415356384
---- n features = 6	MSE: 43.48949924401955	MAE: 5.261063446523245	R^2: 0.341444284453185
---- n features = 7	MSE: 35.32918422873551	MAE: 4.657021901910687	R^2: 0.4696404170024023
---- n features = 8	MSE: 51.0869966225152	MAE: 5.8594843935241965	R^2: 0.22337602969549425

```
---- n features = 9
MSE: 58.94318238498647   MAE: 6.436163727761674   R^2: 0.10728658219899936
Optimal number of features: 7
```

```
[85]: fig, axs = plt.subplots(1, 3, tight_layout=True)
      axs[0].plot(n_feats, mse_nfeat)
      axs[0].set_xlabel("features")
      axs[0].set_ylabel("MSE")

      axs[1].plot(n_feats, mae_nfeat)
      axs[1].set_xlabel("features")
      axs[1].set_ylabel("MAE")

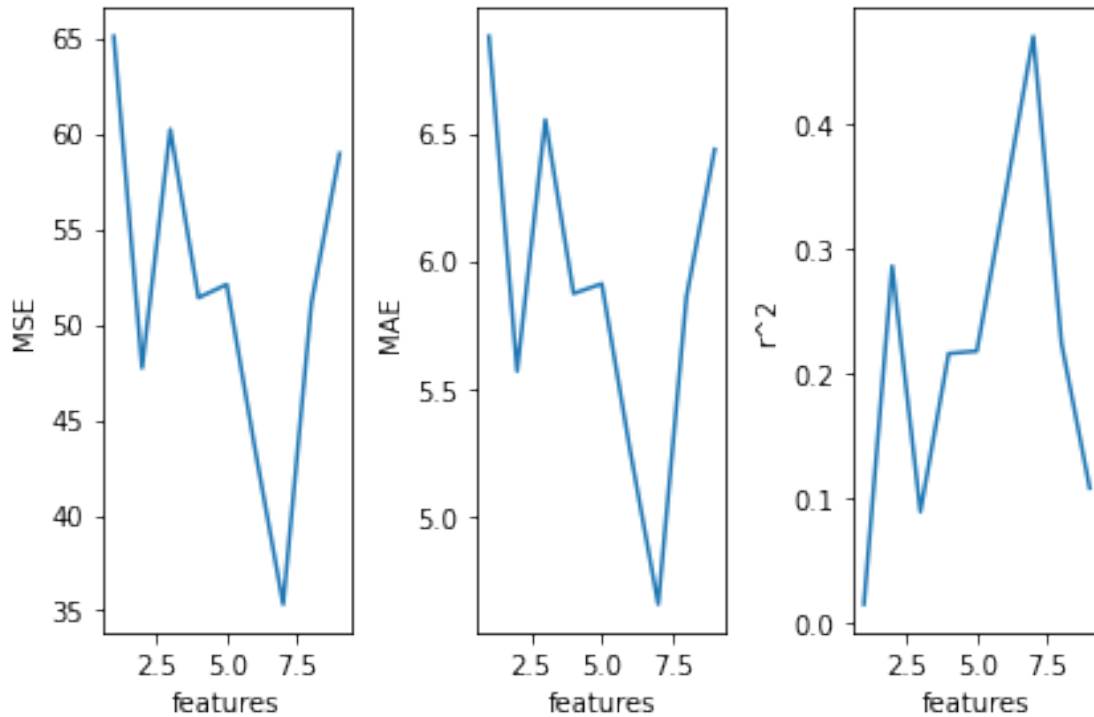
      axs[2].plot(n_feats, r2_nfeat)
      axs[2].set_xlabel("features")
      axs[2].set_ylabel("r^2")

      plt.show()

      # Fit model with optimal number of features
      regr = GradientBoostingRegressor()
      fselection = SequentialFeatureSelector(regr, n_features_to_select =
      ↪opt_features)
      fselection.fit(x, y)

      print("Selected features: ", fselection.get_feature_names_out())

      x_transformed = fselection.transform(x)
      regr.fit(x_transformed, y)
      print("Feature importances:", regr.feature_importances_)
```

Selected features: ['x2' 'x3' 'x5' 'x7' 'x9' 'x13' 'x14']
 Feature importances: [0.10913373 0.03841733 0.15137298 0.06361002 0.13343756
 0.29311274
 0.21091564]

C. Haz el mismo proceso del paso 2, pero ahora con el método de selección de características recursivo. Reporta los predictores óptimos encontrados por el método.

```
[86]: #-----
# Recursive feature selection
#-----

from sklearn.feature_selection import RFE

#####
# Find optimal number of features using cross-validation
#####
print("----- Optimal selection of number of features -----")
print("----- RECURSIVE SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
```

```

r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)

    for train_index, test_index in kf.split(x):

        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]

        regr_cv = GradientBoostingRegressor()

        fselection_cv = RFE(regr_cv, n_features_to_select=n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)

        regr_cv.fit(x_train, y_train)

        # Test phase
        x_test = fselection_cv.transform(x[test_index, :])
        y_test = y[test_index]
        y_pred = regr_cv.predict(x_test)

        mse_i = mean_squared_error(y_test, y_pred)
        mse_cv.append(mse_i)

        mae_i = mean_absolute_error(y_test, y_pred)
        mae_cv.append(mae_i)

        r2_i = r2_score(y_test, y_pred)
        r2_cv.append(r2_i)

    mse = np.average(mse_cv)
    mse_nfeat.append(mse)

    mae = np.average(mae_cv)
    mae_nfeat.append(mae)

    r2 = np.average(r2_cv)
    r2_nfeat.append(r2)

```

```

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

----- Optimal selection of number of features -----
----- RECURSIVE SELECTION -----
---- n features = 1
MSE: 24.86461783840975    MAE: 3.9970016255409795    R^2: 0.6230558429101098
---- n features = 2
MSE: 18.465103118442265    MAE: 3.3032878907294694    R^2: 0.7201387378918725
---- n features = 3
MSE: 15.678117279182391    MAE: 3.075744109062968    R^2: 0.7627149926183354
---- n features = 4
MSE: 14.347686017122225    MAE: 2.9470768819073236    R^2: 0.7826862951211564
---- n features = 5
MSE: 15.229555959345944    MAE: 3.0747640093590016    R^2: 0.7692279324860825
---- n features = 6
MSE: 16.09459340648603    MAE: 3.1433803260398614    R^2: 0.7563658510368928
---- n features = 7
MSE: 15.49263137471694    MAE: 3.111085318772643    R^2: 0.7654670631137804
---- n features = 8
MSE: 15.717534130589996    MAE: 3.135059725881496    R^2: 0.7623853829488703
---- n features = 9
MSE: 15.396241613361013    MAE: 3.106387441820556    R^2: 0.7668538731416393
Optimal number of features: 4

```

```

[88]: fig, axs = plt.subplots(1, 3, tight_layout=True)
      axs[0].plot(n_feats, mse_nfeat)
      axs[0].set_xlabel("features")
      axs[0].set_ylabel("MSE")

      axs[1].plot(n_feats, mae_nfeat)
      axs[1].set_xlabel("features")
      axs[1].set_ylabel("MAE")

      axs[2].plot(n_feats, r2_nfeat)
      axs[2].set_xlabel("features")
      axs[2].set_ylabel("r^2")

      plt.show()

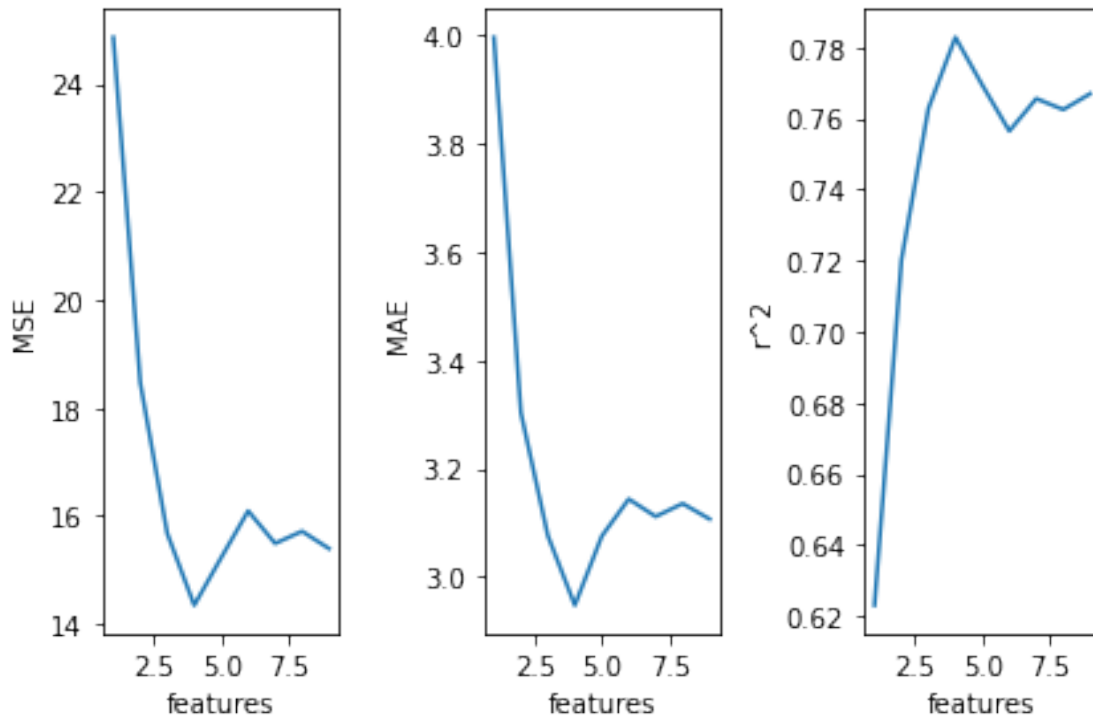
      # Fit model with optimal number of features
      regr = GradientBoostingRegressor()
      fselection = RFE(regr, n_features_to_select = opt_features)

```

```
fselection.fit(x, y)

print("Selected features: ", fselection.get_feature_names_out())

x_transformed = fselection.transform(x)
regr.fit(x_transformed, y)
```



Selected features: ['x0' 'x1' 'x12' 'x14']

[88]: GradientBoostingRegressor()

MODELO 2. DECISION TREE

A. Encuentra el número óptimo de predictores para el modelo utilizando el método filter y validación cruzada. Una vez que tengas el número óptimo, muestra las características seleccionadas.

```
[89]: from sklearn.tree import DecisionTreeRegressor

#-----
# Find optimal number of features using cross-validation
# FILTER
#-----
print("----- Optimal selection of number of features -----")
```

```

print("----- FILTER SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)

    for train_index, test_index in kf.split(x):

        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]

        fselection_cv = SelectKBest(r_regression, k = n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)

        regr_cv = DecisionTreeRegressor()
        regr_cv.fit(x_train, y_train)

        # Test phase
        x_test = fselection_cv.transform(x[test_index, :])
        y_test = y[test_index]
        y_pred = regr_cv.predict(x_test)

        mse_i = mean_squared_error(y_test, y_pred)
        mse_cv.append(mse_i)

        mae_i = mean_absolute_error(y_test, y_pred)
        mae_cv.append(mae_i)

        r2_i = r2_score(y_test, y_pred)
        r2_cv.append(r2_i)

    mse = np.average(mse_cv)
    mse_nfeat.append(mse)

```

```

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

----- Optimal selection of number of features -----
----- FILTER SELECTION -----
---- n features = 1
MSE: 24.330420850375376    MAE: 3.863492599944897    R^2: 0.6313535256803677
---- n features = 2
MSE: 41.61434815517255    MAE: 4.201415638297872    R^2: 0.3696072121648688
---- n features = 3
MSE: 37.26871898959149    MAE: 3.8749996765957446    R^2: 0.43376688409820163
---- n features = 4
MSE: 34.81969746594043    MAE: 3.7289675234042554    R^2: 0.47231571940119677
---- n features = 5
MSE: 32.24197208848    MAE: 3.5658265191489362    R^2: 0.5118185857963713
---- n features = 6
MSE: 32.42228828660085    MAE: 3.5995485276595742    R^2: 0.5077189818875086
---- n features = 7
MSE: 30.113965526180426    MAE: 3.443173872340426    R^2: 0.5436686228282838
---- n features = 8
MSE: 27.282786447537024    MAE: 3.076455285106383    R^2: 0.5871620954796561
---- n features = 9
MSE: 26.657869632302983    MAE: 2.978513157446809    R^2: 0.5956700995300597
Optimal number of features:  1

```

```

[90]: fig, axs = plt.subplots(1, 3, tight_layout=True)
axs[0].plot(n_feats, mse_nfeat)
axs[0].set_xlabel("features")
axs[0].set_ylabel("MSE")

axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("features")
axs[1].set_ylabel("MAE")

axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("features")
axs[2].set_ylabel("r^2")

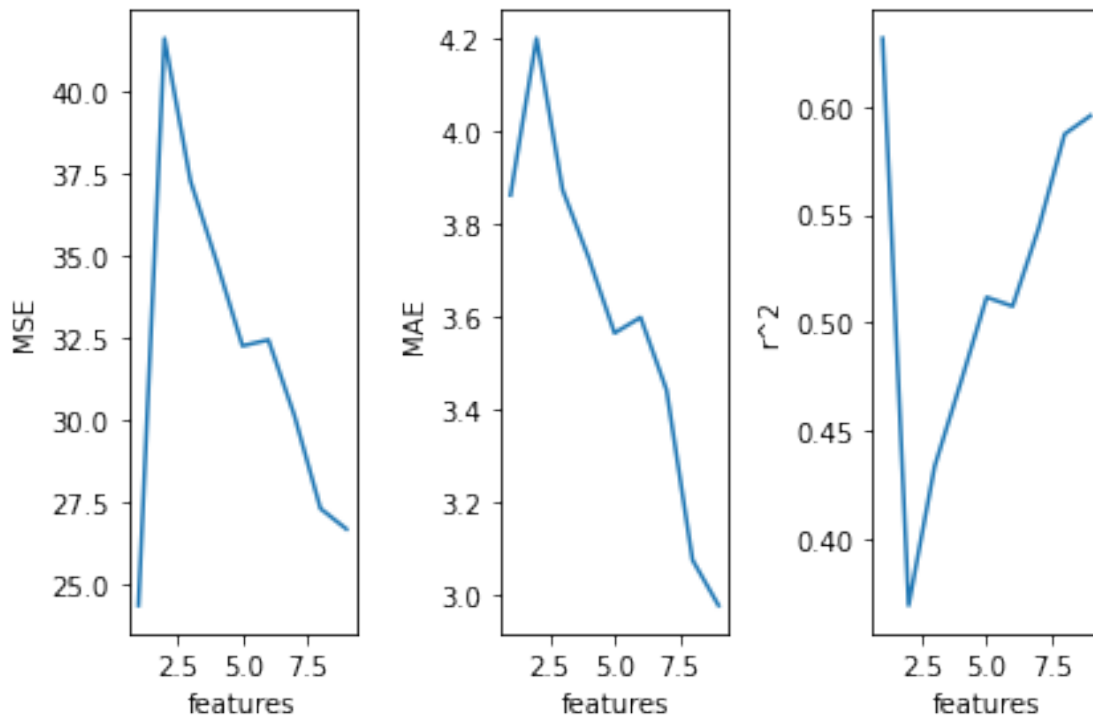
```

```
plt.show()

regr = DecisionTreeRegressor()
fselection = SelectKBest(r_regression, k=opt_features)
fselection.fit(x, y)

print("Selected features: ", fselection.get_feature_names_out())

x_transformed = fselection.transform(x)
regr.fit(x_transformed, y)
```



Selected features: ['x0']

[90]: DecisionTreeRegressor()

B. Repite el paso anterior pero con selección de características secuencial (Wrapper). Reporta los predictores óptimos encontrados por el método.

```
[91]: #####
# Find optimal number of features using cross-validation
#####
print("----- Optimal selection of number of features -----")
print("----- WRAPPER SELECTION -----")
```

```

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)

    for train_index, test_index in kf.split(x):

        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]

        regr_cv = DecisionTreeRegressor()

        fselection_cv = SequentialFeatureSelector(regr_cv,
↪n_features_to_select=n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)

        regr_cv.fit(x_train, y_train)

        # Test phase
        x_test = fselection_cv.transform(x[test_index, :])
        y_test = y[test_index]
        y_pred = regr_cv.predict(x_test)

        mse_i = mean_squared_error(y_test, y_pred)
        mse_cv.append(mse_i)

        mae_i = mean_absolute_error(y_test, y_pred)
        mae_cv.append(mae_i)

        r2_i = r2_score(y_test, y_pred)
        r2_cv.append(r2_i)

    mse = np.average(mse_cv)
    mse_nfeat.append(mse)

```



```

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

----- Optimal selection of number of features -----
----- WRAPPER SELECTION -----
---- n features = 1
MSE: 66.0435246173415    MAE: 6.95674029026232    R^2: -1.6831348966150907e-05
---- n features = 2
MSE: 10.930128687285782    MAE: 2.4818164423520654    R^2: 0.8343313231255316
---- n features = 3
MSE: 6.625720389363583    MAE: 1.094756875744681    R^2: 0.8999262671745631
---- n features = 4
MSE: 10.358575148396595    MAE: 1.4489514382978725    R^2: 0.843430573722389
---- n features = 5
MSE: 9.953543353276595    MAE: 1.204427029787234    R^2: 0.849287478576219
---- n features = 6
MSE: 11.046356726987236    MAE: 1.4959353021276596    R^2: 0.8309735505118965
---- n features = 7
MSE: 11.09206539586383    MAE: 1.526766025531915    R^2: 0.8325104783780362
---- n features = 8
MSE: 8.779937491889363    MAE: 1.1171091234042554    R^2: 0.8671187554287613
---- n features = 9
MSE: 7.033395289068936    MAE: 0.7024832340425533    R^2: 0.8935370738155684
Optimal number of features: 3

```

```

[93]: fig, axs = plt.subplots(1, 3, tight_layout=True)
      axs[0].plot(n_feats, mse_nfeat)
      axs[0].set_xlabel("features")
      axs[0].set_ylabel("MSE")

      axs[1].plot(n_feats, mae_nfeat)
      axs[1].set_xlabel("features")
      axs[1].set_ylabel("MAE")

      axs[2].plot(n_feats, r2_nfeat)
      axs[2].set_xlabel("features")
      axs[2].set_ylabel("r^2")

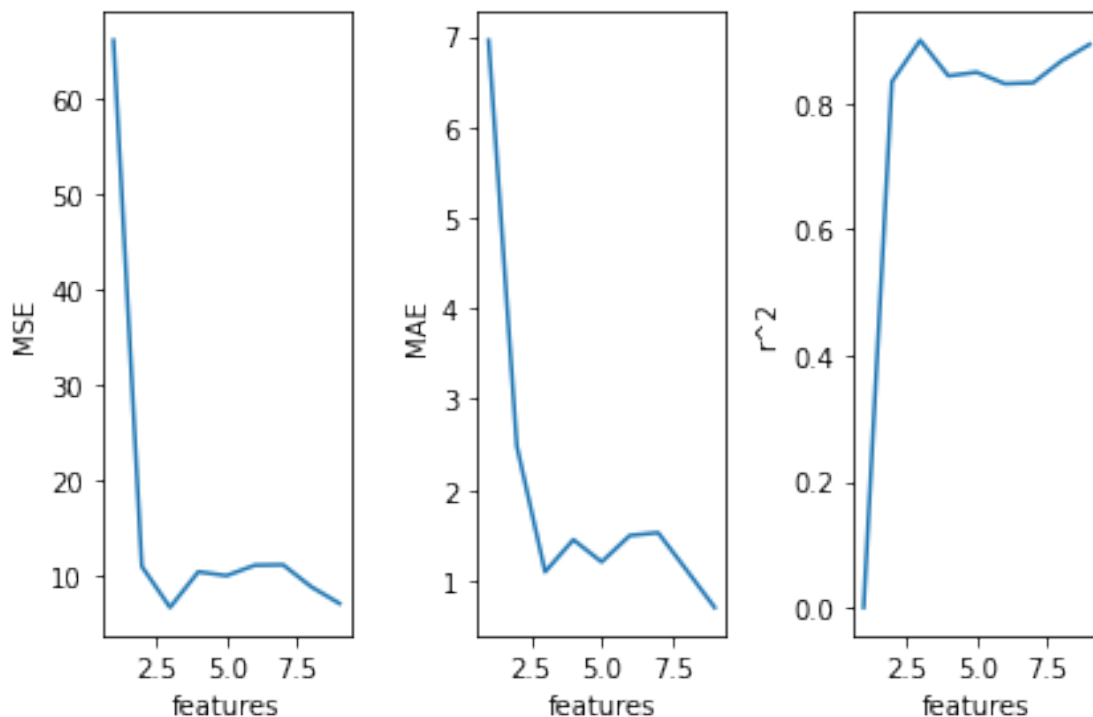
```

```
plt.show()

# Fit model with optimal number of features
regr = DecisionTreeRegressor()
fselection = SequentialFeatureSelector(regr, n_features_to_select = 10,
    ↪ opt_features)
fselection.fit(x, y)

print("Selected features: ", fselection.get_feature_names_out())

x_transformed = fselection.transform(x)
regr.fit(x_transformed, y)
```



Selected features: ['x0' 'x3' 'x14']

[93]: DecisionTreeRegressor()

C. Haz el mismo proceso del paso 2, pero ahora con el método de selección de características recursivo. Reporta los predictores óptimos encontrados por el método.

```
[94]: #-----
# Recursive feature selection
#-----
```

```

from sklearn.feature_selection import RFE

#####
# Find optimal number of features using cross-validation
#####
print("----- Optimal selection of number of features -----")
print("----- RECURSIVE SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)

    for train_index, test_index in kf.split(x):

        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]

        regr_cv = DecisionTreeRegressor()

        fselection_cv = RFE(regr_cv, n_features_to_select=n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)

        regr_cv.fit(x_train, y_train)

        # Test phase
        x_test = fselection_cv.transform(x[test_index, :])
        y_test = y[test_index]
        y_pred = regr_cv.predict(x_test)

        mse_i = mean_squared_error(y_test, y_pred)
        mse_cv.append(mse_i)

        mae_i = mean_absolute_error(y_test, y_pred)
        mae_cv.append(mae_i)

```

```

        r2_i = r2_score(y_test, y_pred)
        r2_cv.append(r2_i)

    mse = np.average(mse_cv)
    mse_nfeat.append(mse)

    mae = np.average(mae_cv)
    mae_nfeat.append(mae)

    r2 = np.average(r2_cv)
    r2_nfeat.append(r2)

    print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

----- Optimal selection of number of features -----
----- RECURSIVE SELECTION -----
---- n features = 1
MSE: 24.317127780094214    MAE: 3.862374951490293    R^2: 0.630755885682157
---- n features = 2
MSE: 10.835329904726908    MAE: 1.3326313089361703    R^2: 0.8376312975373013
---- n features = 3
MSE: 11.033082337112166    MAE: 0.9374294164539008    R^2: 0.8335359368318095
---- n features = 4
MSE: 3.5225638814195745    MAE: 0.33721879148936174    R^2: 0.9467964398790268
---- n features = 5
MSE: 3.3661776655387228    MAE: 0.35412847659574476    R^2: 0.9487533217559445
---- n features = 6
MSE: 4.003043076999149    MAE: 0.438547455319149    R^2: 0.939602012132086
---- n features = 7
MSE: 4.7834007687455316    MAE: 0.5260925617021277    R^2: 0.9274233110773524
---- n features = 8
MSE: 5.236698277712341    MAE: 0.5596859914893618    R^2: 0.9206673491932087
---- n features = 9
MSE: 5.702739787911489    MAE: 0.5924543659574468    R^2: 0.9137570038493517
Optimal number of features: 5

```

```

[96]: fig, axs = plt.subplots(1, 3, tight_layout=True)
      axs[0].plot(n_feats, mse_nfeat)
      axs[0].set_xlabel("features")
      axs[0].set_ylabel("MSE")

      axs[1].plot(n_feats, mae_nfeat)

```

```

axs[1].set_xlabel("features")
axs[1].set_ylabel("MAE")

axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("features")
axs[2].set_ylabel("r^2")

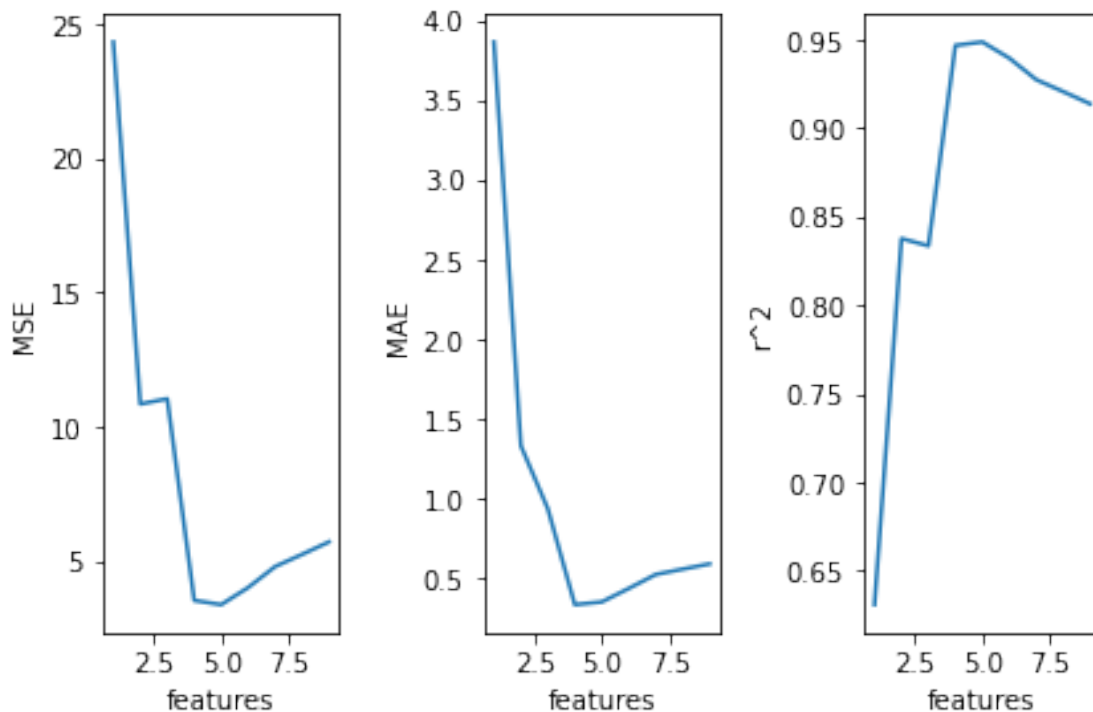
plt.show()

# Fit model with optimal number of features
regr = DecisionTreeRegressor()
fselection = RFE(regr, n_features_to_select = opt_features)
fselection.fit(x, y)

print("Selected features: ", fselection.get_feature_names_out())

x_transformed = fselection.transform(x)
regr.fit(x_transformed, y)

```



Selected features: ['x0' 'x1' 'x7' 'x12' 'x14']

[96]: DecisionTreeRegressor()

MODELO 3. RANDOM FOREST

A. Encuentra el número óptimo de predictores para el modelo utilizando el método filter y validación cruzada. Una vez que tengas el número óptimo, muestra las características seleccionadas.

```
[97]: from sklearn.ensemble import RandomForestRegressor

#-----
# Find optimal number of features using cross-validation
# FILTER
#-----

print("----- Optimal selection of number of features -----")
print("----- FILTER SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)

    for train_index, test_index in kf.split(x):

        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]

        fselection_cv = SelectKBest(r_regression, k = n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)

        regr_cv = RandomForestRegressor()
        regr_cv.fit(x_train, y_train)

        # Test phase
        x_test = fselection_cv.transform(x[test_index, :])
        y_test = y[test_index]
        y_pred = regr_cv.predict(x_test)

        mse_i = mean_squared_error(y_test, y_pred)
        mse_cv.append(mse_i)
```

```

    mae_i = mean_absolute_error(y_test, y_pred)
    mae_cv.append(mae_i)

    r2_i = r2_score(y_test, y_pred)
    r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

----- Optimal selection of number of features -----
----- FILTER SELECTION -----
---- n features = 1
MSE: 24.369287940976626    MAE: 3.8654491894711773    R^2: 0.6311559063715928
---- n features = 2
MSE: 29.296128019348714    MAE: 3.697492356290037    R^2: 0.5559587108581364
---- n features = 3
MSE: 22.232171269165157    MAE: 3.2476887376235055    R^2: 0.6629787529070851
---- n features = 4
MSE: 19.651860545757387    MAE: 3.0719339012765956    R^2: 0.7019853124717809
---- n features = 5
MSE: 19.2151816090298    MAE: 3.0381923002553193    R^2: 0.7083140863666524
---- n features = 6
MSE: 18.093362562555008    MAE: 2.9689962873191487    R^2: 0.724677712617431
---- n features = 7
MSE: 17.53503610317096    MAE: 2.913954605617021    R^2: 0.7340894829641093
---- n features = 8
MSE: 15.683308487624876    MAE: 2.765015600510638    R^2: 0.7625409732491084
---- n features = 9
MSE: 13.604076249945077    MAE: 2.4519290982127657    R^2: 0.7940875456167605
Optimal number of features: 9

```

```

[98]: fig, axs = plt.subplots(1, 3, tight_layout=True)
      axs[0].plot(n_feats, mse_nfeat)
      axs[0].set_xlabel("k")

```

```

axs[0].set_ylabel("MSE")

axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("k")
axs[1].set_ylabel("MAE")

axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("k")
axs[2].set_ylabel("r^2")

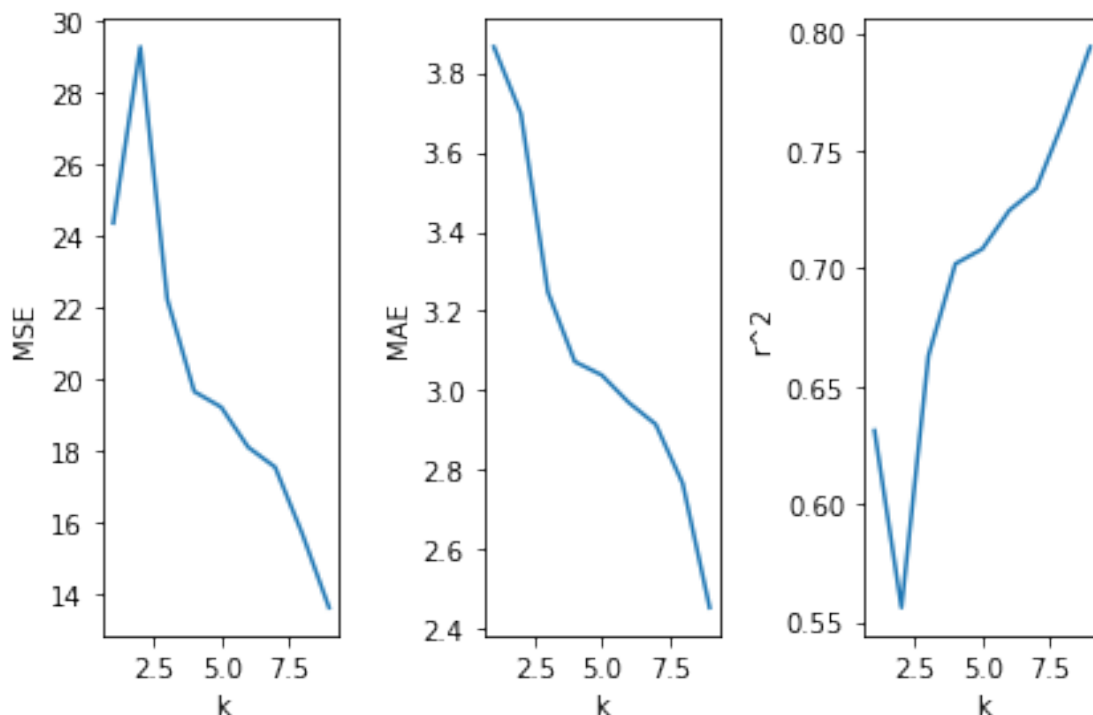
plt.show()

# Fit model with optimal number of features
regr = RandomForestRegressor()
fselection = SelectKBest(r_regression, k = opt_features)
fselection.fit(x, y)

print("Selected features: ", fselection.get_feature_names_out())

x_transformed = fselection.transform(x)
regr.fit(x_transformed, y)
print("Feature importances:", regr.feature_importances_)

```



Selected features: ['x0' 'x2' 'x4' 'x6' 'x7' 'x8' 'x9' 'x10' 'x13']

Feature importances: [0.66544806 0.04137571 0.05891652 0.03013702 0.02543858
0.04321219
0.02830603 0.05919288 0.04797301]

**B. Repite el paso anterior pero con selección de características secuencial (Wrapper).
Reporta los predictores óptimos encontrados por el método.**

```
[99]: #####  
# Find optimal number of features using cross-validation  
#####  
print("----- Optimal selection of number of features -----")  
print("----- WRAPPER SELECTION -----")  
  
n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]  
  
mse_nfeat = []  
mae_nfeat = []  
r2_nfeat = []  
for n_feat in n_feats:  
    print('---- n features =', n_feat)  
  
    mse_cv = []  
    mae_cv = []  
    r2_cv = []  
  
    kf = KFold(n_splits=5, shuffle = True)  
  
    for train_index, test_index in kf.split(x):  
  
        # Training phase  
        x_train = x[train_index, :]  
        y_train = y[train_index]  
  
        regr_cv = RandomForestRegressor()  
  
        fselection_cv = SequentialFeatureSelector(regr_cv,   
↪n_features_to_select=n_feat)  
        fselection_cv.fit(x_train, y_train)  
        x_train = fselection_cv.transform(x_train)  
  
        regr_cv.fit(x_train, y_train)  
  
        # Test phase  
        x_test = fselection_cv.transform(x[test_index, :])  
        y_test = y[test_index]  
        y_pred = regr_cv.predict(x_test)  
  
        mse_i = mean_squared_error(y_test, y_pred)
```

```

mse_cv.append(mse_i)

mae_i = mean_absolute_error(y_test, y_pred)
mae_cv.append(mae_i)

r2_i = r2_score(y_test, y_pred)
r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

```

----- Optimal selection of number of features -----

----- WRAPPER SELECTION -----

n features	MSE	MAE	R ²
1	66.09389822246051	6.960247110669341	-0.0013378153594672648
2	36.387816383503306	4.343825719202871	0.4483800118164057
3	62.9722148148259	6.143004408331509	0.0555607168859545
4	38.56700272112519	4.145245411347518	0.428246222129218
5	37.97460617376798	4.336239997446809	0.41448931337780265
6	38.25378316676401	4.60548174280851	0.41997726927007495
7	45.142255182441616	5.189304875404255	0.3175248386772275
8	32.69346771728972	3.794176868425532	0.5026646228069207
9	43.22934335866652	4.845802464340425	0.3479471610074273

Optimal number of features: 8

```
[100]: fig, axs = plt.subplots(1, 3, tight_layout=True)
axs[0].plot(n_feats, mse_nfeat)
axs[0].set_xlabel("features")
axs[0].set_ylabel("MSE")

axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("features")
axs[1].set_ylabel("MAE")

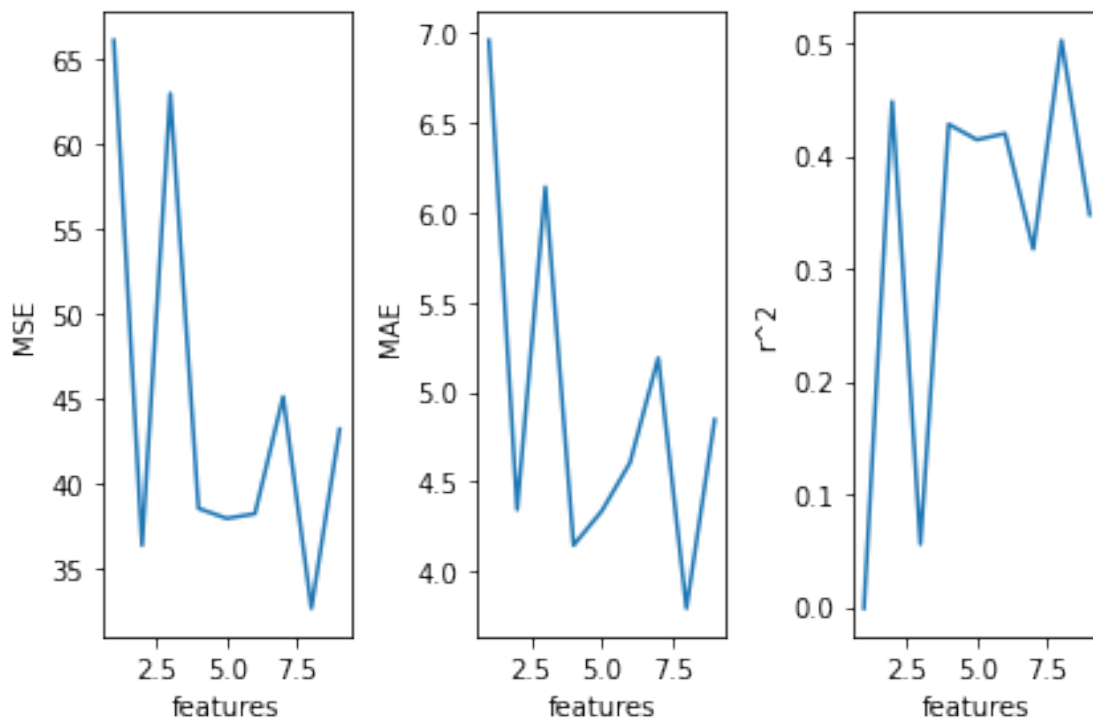
axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("features")
axs[2].set_ylabel("r^2")

plt.show()

regr = RandomForestRegressor()
fselection = SequentialFeatureSelector(regr, n_features_to_select=opt_features)
fselection.fit(x, y)

print("Selected features: ", fselection.get_feature_names_out())

x_transformed = fselection.transform(x)
regr.fit(x_transformed, y)
```



Selected features: ['x1' 'x3' 'x5' 'x6' 'x7' 'x10' 'x13' 'x14']

[100]: RandomForestRegressor()

C. Haz el mismo proceso del paso 2, pero ahora con el método de selección de características recursivo. Reporta los predictores óptimos encontrados por el método.

```
[101]: #-----
# Recursive feature selection
#-----

from sklearn.feature_selection import RFE

#####
# Find optimal number of features using cross-validation
#####
print("----- Optimal selection of number of features -----")
print("----- RECURSIVE SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)

    for train_index, test_index in kf.split(x):

        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]

        regr_cv = RandomForestRegressor()

        fselection_cv = RFE(regr_cv, n_features_to_select=n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)

        regr_cv.fit(x_train, y_train)

    # Test phase
```

```

x_test = fselection_cv.transform(x[test_index, :])
y_test = y[test_index]
y_pred = regr_cv.predict(x_test)

mse_i = mean_squared_error(y_test, y_pred)
mse_cv.append(mse_i)

mae_i = mean_absolute_error(y_test, y_pred)
mae_cv.append(mae_i)

r2_i = r2_score(y_test, y_pred)
r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

----- Optimal selection of number of features -----
----- RECURSIVE SELECTION -----
---- n features = 1
MSE: 24.272947147965418    MAE: 3.8612975331394823    R^2: 0.6322584465394725
---- n features = 2
MSE: 20.473261401955675    MAE: 2.974728264009321    R^2: 0.689916942748059
---- n features = 3
MSE: 7.5278891178122205    MAE: 1.2000291222127664    R^2: 0.8859740320376863
---- n features = 4
MSE: 2.275380504820601    MAE: 0.4518507485957448    R^2: 0.9653019584900221
---- n features = 5
MSE: 2.0007105729771304    MAE: 0.4313618316595743    R^2: 0.9697508985785909
---- n features = 6
MSE: 1.7846592593844697    MAE: 0.4655994410212765    R^2: 0.9731096257871723
---- n features = 7
MSE: 2.029178055503246    MAE: 0.5330757111489363    R^2: 0.9694558677857111
---- n features = 8
MSE: 2.178052245017098    MAE: 0.577453378042553    R^2: 0.9669241346666141
---- n features = 9

```

MSE: 2.242051792313009 MAE: 0.6053375339574467 R²: 0.9660159930187516
Optimal number of features: 6

```
[102]: fig, axs = plt.subplots(1, 3, tight_layout=True)
axs[0].plot(n_feats, mse_nfeat)
axs[0].set_xlabel("features")
axs[0].set_ylabel("MSE")

axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("features")
axs[1].set_ylabel("MAE")

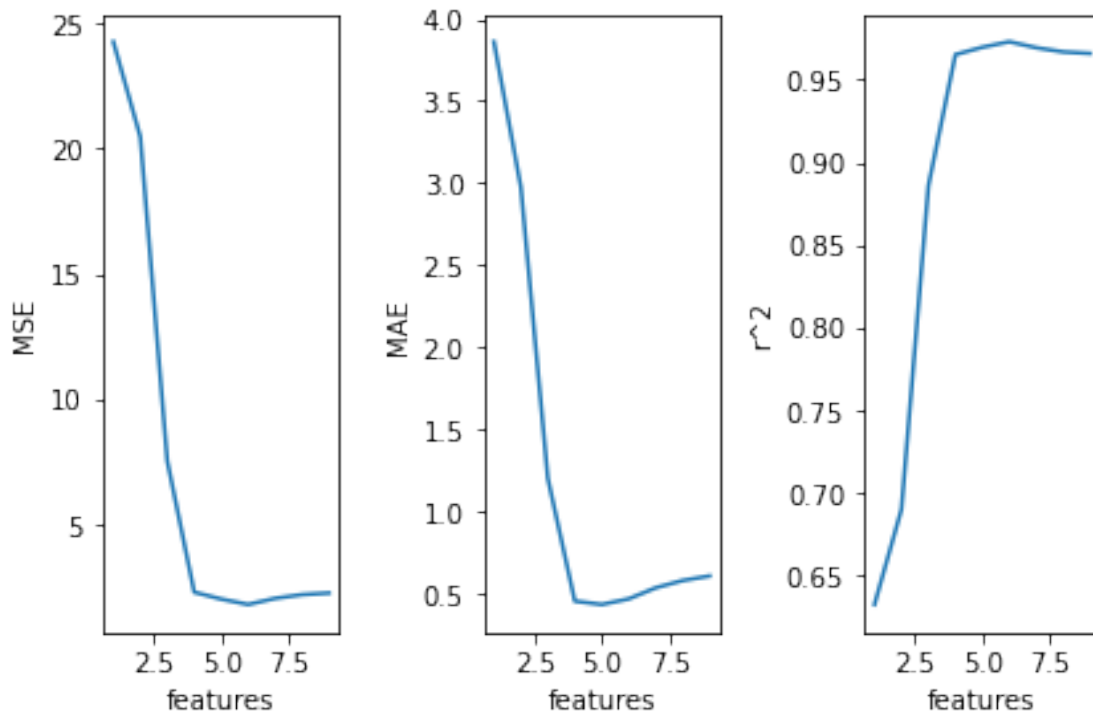
axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("features")
axs[2].set_ylabel("r^2")

plt.show()

# Fit model with optimal number of features
regr = RandomForestRegressor()
fselection = RFE(regr, n_features_to_select = opt_features)
fselection.fit(x, y)

print("Selected features: ", fselection.get_feature_names_out())

x_transformed = fselection.transform(x)
regr.fit(x_transformed, y)
```



Selected features: ['x0' 'x1' 'x9' 'x11' 'x12' 'x14']

[102]: RandomForestRegressor()

MODELO 4. SVR

A. Encuentra el número óptimo de predictores para el modelo utilizando el método filter y validación cruzada. Una vez que tengas el número óptimo, muestra las características seleccionadas.

```
[103]: from sklearn.svm import SVR
#-----
# Find optimal number of features using cross-validation
# FILTER
#-----
print("----- Optimal selection of number of features -----")
print("----- FILTER SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('----- n features =', n_feat)
```

```

mse_cv = []
mae_cv = []
r2_cv = []

kf = KFold(n_splits=5, shuffle = True)

for train_index, test_index in kf.split(x):

    # Training phase
    x_train = x[train_index, :]
    y_train = y[train_index]

    fselection_cv = SelectKBest(r_regression, k = n_feat)
    fselection_cv.fit(x_train, y_train)
    x_train = fselection_cv.transform(x_train)

    regr_cv = SVR()
    regr_cv.fit(x_train, y_train)

    # Test phase
    x_test = fselection_cv.transform(x[test_index, :])
    y_test = y[test_index]
    y_pred = regr_cv.predict(x_test)

    mse_i = mean_squared_error(y_test, y_pred)
    mse_cv.append(mse_i)

    mae_i = mean_absolute_error(y_test, y_pred)
    mae_cv.append(mae_i)

    r2_i = r2_score(y_test, y_pred)
    r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)

```



```

opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

----- Optimal selection of number of features -----
----- FILTER SELECTION -----
---- n features = 1
MSE: 54.55848170106956    MAE: 5.8939441236218    R^2: 0.1739772547242894
---- n features = 2
MSE: 60.84106700199023    MAE: 6.6766149225289295    R^2: 0.07883516951607743
---- n features = 3
MSE: 61.09954894771947    MAE: 6.686799119347681    R^2: 0.07462294723054715
---- n features = 4
MSE: 61.08856549404574    MAE: 6.685390405481121    R^2: 0.07526851696123382
---- n features = 5
MSE: 61.156749409114376    MAE: 6.688089273196503    R^2: 0.07259825821275197
---- n features = 6
MSE: 61.27047151238828    MAE: 6.691879872988267    R^2: 0.072001178123722
---- n features = 7
MSE: 61.2288402905268    MAE: 6.689674780594333    R^2: 0.07200896250582281
---- n features = 8
MSE: 61.109070415220245    MAE: 6.685684564260681    R^2: 0.074973437910721
---- n features = 9
MSE: 61.26820254686695    MAE: 6.6875014620258755    R^2: 0.07252110062950365
Optimal number of features:  1

```

```

[105]: fig, axs = plt.subplots(1, 3, tight_layout=True)
axs[0].plot(n_feats, mse_nfeat)
axs[0].set_xlabel("features")
axs[0].set_ylabel("MSE")

axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("features")
axs[1].set_ylabel("MAE")

axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("features")
axs[2].set_ylabel("r^2")

plt.show()

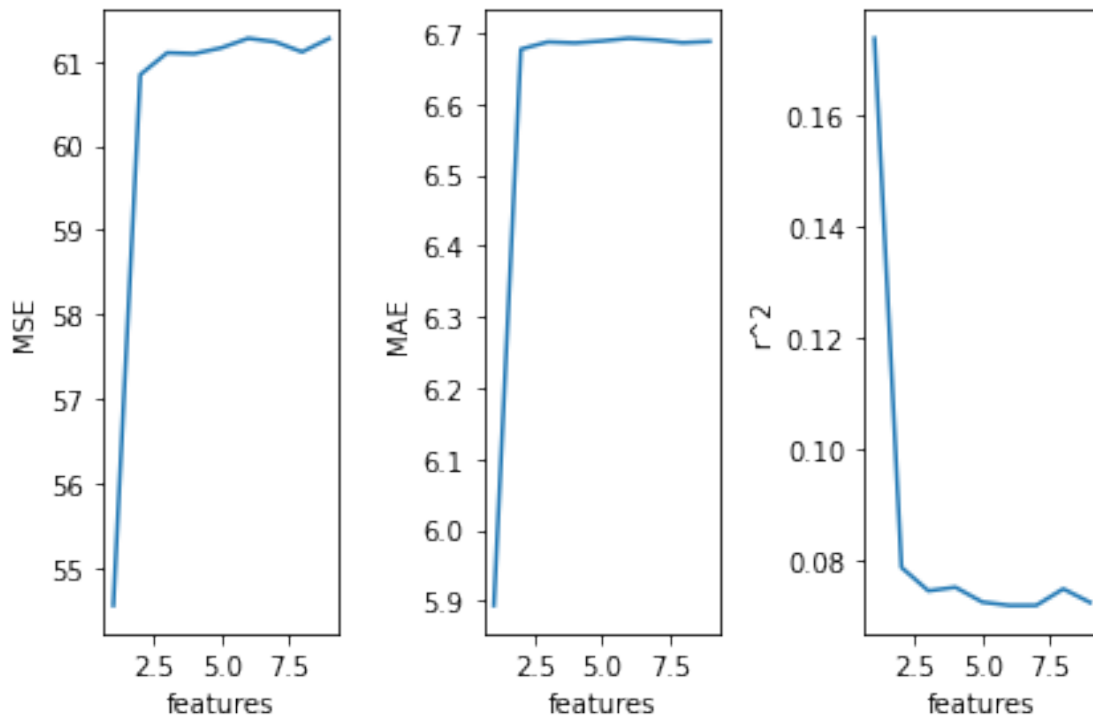
svr = SVR()
fselection = SelectKBest(r_regression, k=opt_features)
fselection.fit(x, y)

print("Selected features: ", fselection.get_feature_names_out())

x_transformed = fselection.transform(x)

```

```
svr.fit(x_transformed, y)
```



Selected features: ['x0']

[105]: SVR()

B. Repite el paso anterior pero con selección de características secuencial (Wrapper).
Reporta los predictores óptimos encontrados por el método.

```
[106]: #####
# Find optimal number of features using cross-validation
#####
print("----- Optimal selection of number of features -----")
print("----- WRAPPER SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
```

```

mae_cv = []
r2_cv = []

kf = KFold(n_splits=5, shuffle = True)

for train_index, test_index in kf.split(x):

    # Training phase
    x_train = x[train_index, :]
    y_train = y[train_index]

    regr_cv = SVR()

    fselection_cv = SequentialFeatureSelector(regr_cv,
↪n_features_to_select=n_feat)
    fselection_cv.fit(x_train, y_train)
    x_train = fselection_cv.transform(x_train)

    regr_cv.fit(x_train, y_train)

    # Test phase
    x_test = fselection_cv.transform(x[test_index, :])
    y_test = y[test_index]
    y_pred = regr_cv.predict(x_test)

    mse_i = mean_squared_error(y_test, y_pred)
    mse_cv.append(mse_i)

    mae_i = mean_absolute_error(y_test, y_pred)
    mae_cv.append(mae_i)

    r2_i = r2_score(y_test, y_pred)
    r2_cv.append(r2_i)

mse = np.average(mse_cv)
mse_nfeat.append(mse)

mae = np.average(mae_cv)
mae_nfeat.append(mae)

r2 = np.average(r2_cv)
r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)

```

```

opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

----- Optimal selection of number of features -----
----- WRAPPER SELECTION -----
---- n features = 1
MSE: 65.33764680454135    MAE: 6.84779625158914    R^2: 0.010162964211005043
---- n features = 2
MSE: 64.81434665818774    MAE: 6.801856685707423    R^2: 0.018879853470512353
---- n features = 3
MSE: 64.54441123897865    MAE: 6.7382849151238915    R^2: 0.02247463734313433
---- n features = 4
MSE: 64.14090896973114    MAE: 6.725450183532317    R^2: 0.02900143252407781
---- n features = 5
MSE: 64.10253103457156    MAE: 6.716609362308129    R^2: 0.029515330263070228
---- n features = 6
MSE: 62.71876828109106    MAE: 6.6497592507415515    R^2: 0.05049283515407497
---- n features = 7
MSE: 63.59044229981955    MAE: 6.698153757598851    R^2: 0.03695565327797723
---- n features = 8
MSE: 63.94671742998577    MAE: 6.727007334761234    R^2: 0.03030913840149201
---- n features = 9
MSE: 62.07699191355846    MAE: 6.648801622329051    R^2: 0.06039446483749302
Optimal number of features: 9

```

```

[107]: fig, axs = plt.subplots(1, 3, tight_layout=True)
axs[0].plot(n_feats, mse_nfeat)
axs[0].set_xlabel("features")
axs[0].set_ylabel("MSE")

axs[1].plot(n_feats, mae_nfeat)
axs[1].set_xlabel("features")
axs[1].set_ylabel("MAE")

axs[2].plot(n_feats, r2_nfeat)
axs[2].set_xlabel("features")
axs[2].set_ylabel("r^2")

plt.show()

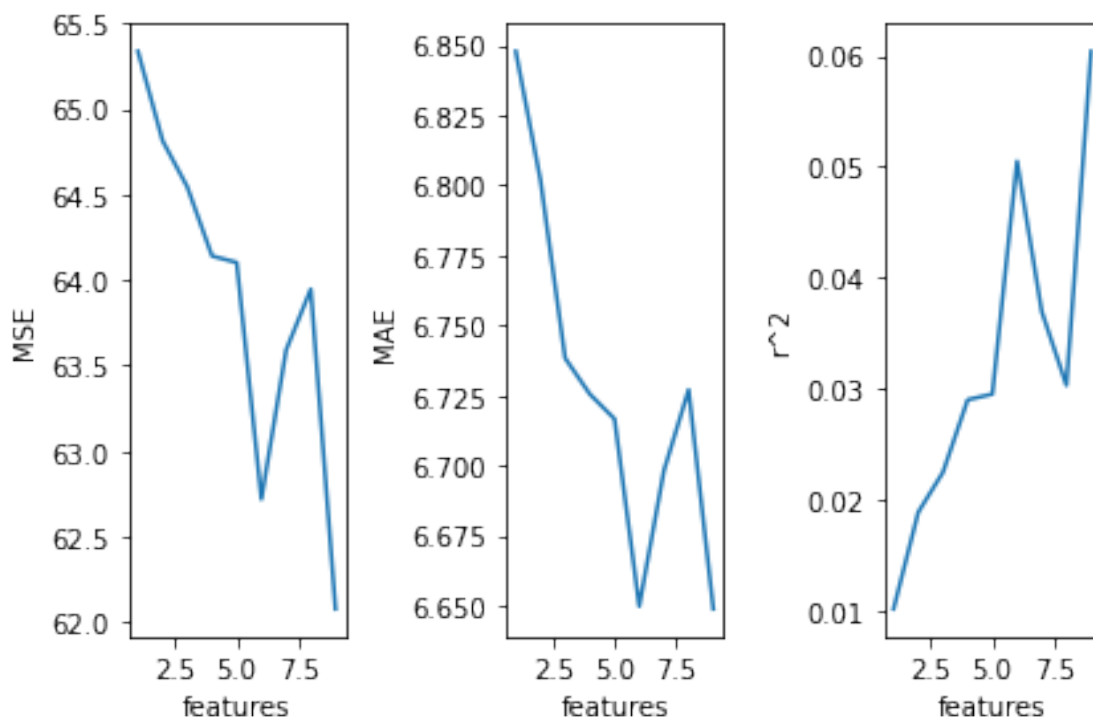
svr = SVR()
fselection = SequentialFeatureSelector(svr, n_features_to_select=opt_features)
fselection.fit(x, y)

print("Selected features: ", fselection.get_feature_names_out())

x_transformed = fselection.transform(x)

```

```
svr.fit(x_transformed, y)
```



Selected features: ['x2' 'x3' 'x4' 'x5' 'x7' 'x8' 'x10' 'x13' 'x14']

[107]: SVR()

C. Haz el mismo proceso del paso 2, pero ahora con el método de selección de características recursivo. Reporta los predictores óptimos encontrados por el método.

```
[14]: #-----
# Recursive feature selection
#-----
from sklearn.svm import LinearSVR
from sklearn.feature_selection import RFE

#####
# Find optimal number of features using cross-validation
#####
print("----- Optimal selection of number of features -----")
print("----- RECURSIVE SELECTION -----")

n_feats = [1, 2, 3, 4, 5, 6, 7, 8, 9]

mse_nfeat = []
```

```

mae_nfeat = []
r2_nfeat = []
for n_feat in n_feats:
    print('---- n features =', n_feat)

    mse_cv = []
    mae_cv = []
    r2_cv = []

    kf = KFold(n_splits=5, shuffle = True)

    for train_index, test_index in kf.split(x):

        # Training phase
        x_train = x[train_index, :]
        y_train = y[train_index]

        regr_cv = LinearSVR()

        fselection_cv = RFE(regr_cv, n_features_to_select=n_feat)
        fselection_cv.fit(x_train, y_train)
        x_train = fselection_cv.transform(x_train)

        regr_cv.fit(x_train, y_train)

        # Test phase
        x_test = fselection_cv.transform(x[test_index, :])
        y_test = y[test_index]
        y_pred = regr_cv.predict(x_test)

        mse_i = mean_squared_error(y_test, y_pred)
        mse_cv.append(mse_i)

        mae_i = mean_absolute_error(y_test, y_pred)
        mae_cv.append(mae_i)

        r2_i = r2_score(y_test, y_pred)
        r2_cv.append(r2_i)

    mse = np.average(mse_cv)
    mae_nfeat.append(mse)

    mae = np.average(mae_cv)
    mae_nfeat.append(mae)

    r2 = np.average(r2_cv)

```

```

r2_nfeat.append(r2)

print('MSE:', mse, ' MAE:', mae, ' R^2:', r2)

opt_index = np.argmin(mse_nfeat)
opt_features = n_feats[opt_index]
print("Optimal number of features: ", opt_features)

----- Optimal selection of number of features -----
----- RECURSIVE SELECTION -----
---- n features = 1
MSE: 65.78625083407674    MAE: 6.877335962308193    R^2: 0.0028457474876315337
---- n features = 2
MSE: 63.16583746055642    MAE: 6.749456933776592    R^2: 0.04317326192120559
---- n features = 3
MSE: 63.13365214338027    MAE: 6.742140017971579    R^2: 0.04432969840352916
---- n features = 4
MSE: 63.26593310370892    MAE: 6.733913406392304    R^2: 0.04210578659009816
---- n features = 5
MSE: 63.125848841860964    MAE: 6.718169384992359    R^2: 0.04426431467423737
---- n features = 6
MSE: 63.3695520295385    MAE: 6.688718398068265    R^2: 0.040755260981329956
---- n features = 7
MSE: 63.574862382021436    MAE: 6.698921953507842    R^2: 0.03725372595051331
---- n features = 8
MSE: 63.3527892931336    MAE: 6.669598776282157    R^2: 0.04031579193023542
---- n features = 9
MSE: 64.16453370449692    MAE: 6.619504112310385    R^2: 0.02844772428789799
Optimal number of features: 5

```

```

[17]: fig, axs = plt.subplots(1, 3, tight_layout=True)
      axs[0].plot(n_feats, mse_nfeat)
      axs[0].set_xlabel("features")
      axs[0].set_ylabel("MSE")

      axs[1].plot(n_feats, mae_nfeat)
      axs[1].set_xlabel("features")
      axs[1].set_ylabel("MAE")

      axs[2].plot(n_feats, r2_nfeat)
      axs[2].set_xlabel("features")
      axs[2].set_ylabel("r^2")

      plt.show()

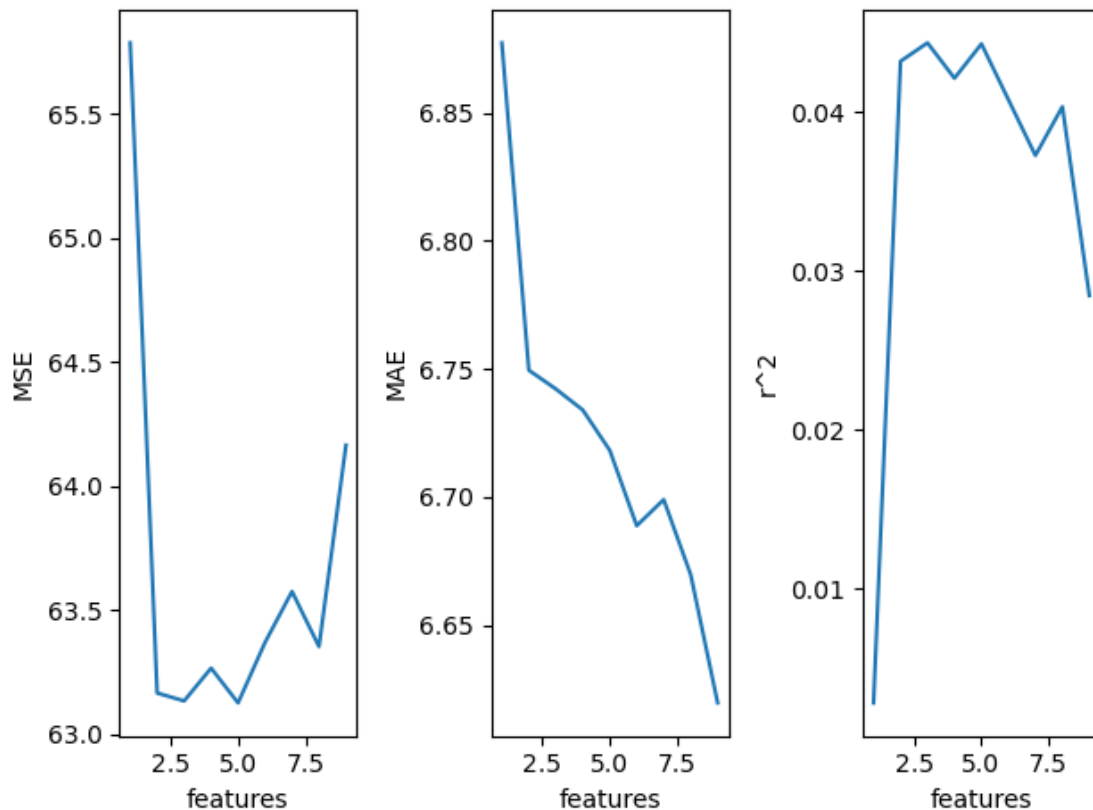
      svr = LinearSVR()
      fselection = RFE(svr, n_features_to_select=opt_features)

```

```
fselection.fit(x, y)

print("Selected features: ", fselection.support_)

x_transformed = fselection.transform(x)
svr.fit(x_transformed, y)
```



```
Selected features: [False False False False False False  True False False  True
                  True False
                  True  True False]
```

```
[17]: LinearSVR()
```

7. Viendo los resultados de este ejercicio, escriba una conclusión sobre los siguientes puntos:

(a) Consideras que el modelo de regresión lineal es adecuado para los datos. ¿Por qué?

Definitivamente no es el modelo lineal el más adecuado ya que el R^2 es 0.0930 y los valores de errores MSE y MAE son muy grandes.

(b) ¿Qué método de selección de características consideras que funciona bien con los datos? ¿Por qué?

Random Forest porque tuvo menores errores y mejor valor de R^2 MSE: 2.2420 MAE: 0.605337 R^2 : 0.96601

(c) Del proceso de selección de características, ¿puedes identificar algunas que sean sobresalientes?

Si. Las características sobresalientes son 'x0' 'x1' 'x9' 'x11' 'x12' 'x14', osea Adult mortality, Polio, HIV/AIDS , Population.

(d) ¿Los modelos de regresión no lineal funcionaron mejor que el lineal? ¿Por qué?

Si fueron mejores los modelos no lineales porque el valor de los errores fue menor y con un R^2 mayor, es decir, se ajustaban mejor a los datos.

(e) ¿Se puede concluir algo interesante sobre los resultados de modelar estos datos con regresión?

Si, que las relaciones entre las variables no son lineales, lo que implica que se requieren de modelos y técnicas complejas.